

# Azure Data Warehouse In-A-Day

Steve Young - Data & AI CSA  
Rebecca Young - Data & AI CSA



# Agenda

## Agenda:

Time	Topic	Description	Materials
09:00am - 09:15am	Introductions & Logistics (15min)	Welcome	N/A
09:15am - 10:00am	Datawarehouse Patterns in Azure & SQL DW Overview (45min)	Slide Deck 01	N/A
10:00am - 10:45am	SQL DW Gen2 New Features & Planning Your Project Build (45min)	Slide Deck 02	N/A
10:45am - 11:00pm	Break (15min)	Please take a break	N/A
11:00am - 12:00pm	Demo & Lab 01 (60 Min)	Setting up the LAB environment	Lab 01
12:00pm - 1:00pm	Lunch (60 Min)	Lunch and complete lab 01	N/A
01:00pm - 1:30pm	SQLDW Loading Best Practices (30 Min)	Lecture	N/A
01:30pm - 02:15pm	Lab 02/03: User IDs & Data loading scenarios and best practices (45min)	Loading different scenarios	Lab 02/03
02:15am - 2:30pm	Break (15min)	Please take a break	N/A
02:30pm - 3:00pm	SQLDW Operational Best Practices (30 Min)	Lecture	N/A
03:00pm - 03:45pm	Lab 04: Performance Tuning best practices (45min)		Lab 04
03:45pm - 4:15pm	Lab 05: Lab 3: Monitoring, Maintenance and Security (30min)		Lab 05
4:15pm - 5:00pm	Q&A and Wrap-up (45min)	final remarks or takeaways/next steps	Survey

# SQL DW Operational best practices

# Topics

- Data Movement
- Replicated Tables
- Concurrency, Concurrency Slots and Resource Classes
- Index/Stats
- Monitoring


# Data and Compute are Distributed

- Massively Parallel Processing (MPP)
  - Multiple compute nodes with dedicated CPU, memory, storage
  - Handles and **hides query complexity**
- Good for scalability
  - Data is spread (distributed) across servers
- Introduces overheads
  - Data is not all in the same place!
  - Don't know where specific values are located
  - **May need to move the data** then process it

SMP Query Execution



MPP Query Execution



# Data Movement



# Why Data Moves

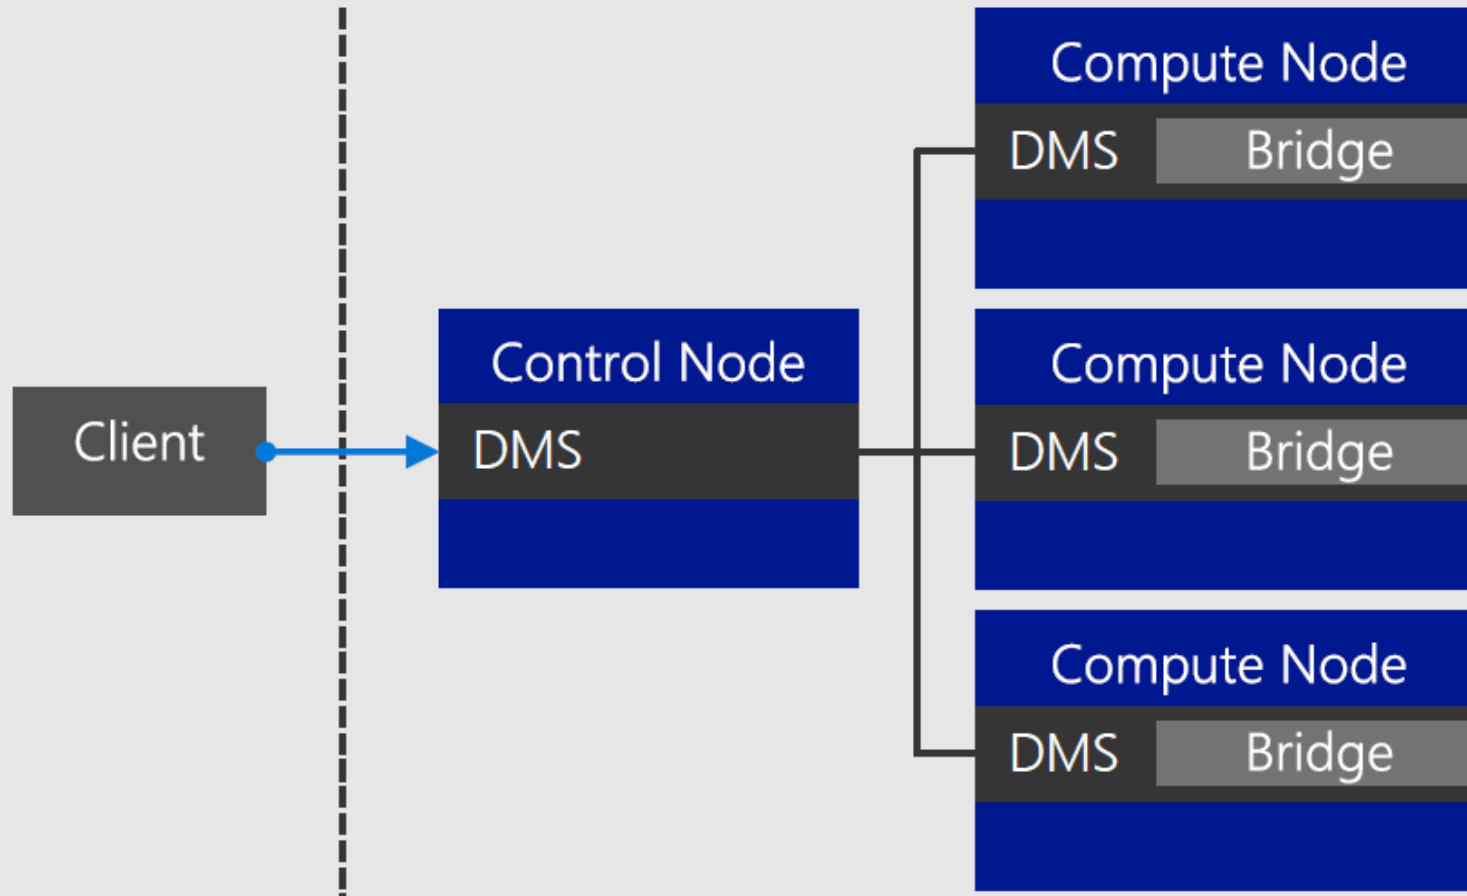
Data has to be co-located to be operated on...

Common reasons:

- Incompatible join

- Incompatible aggregation

# Why Data Moves





# Data Movement Types for a Query

DMS Operation	Description
ShuffleMoveOperation	Distribution → Hash algorithm → New distribution Changing the distribution column in preparation for join.
PartitionMoveOperation	Distribution → Control Node Aggregations - count(*) is count on nodes, sum of count
BroadcastMoveOperation	Distribution → Copy to all distributions Changes distributed table to replicated table for join.
TrimMoveOperation	Replicated table → Hash algorithm → Distribution When a replicated table needs to become distributed. Needed for outer joins.
MoveOperation	Control Node → Copy to all distributions Data moved from Control Node back to Compute Nodes resulting in a replicated table for further processing.


# Hash Distributed Data Movement

## Hash Distributed

ProductSales – Raw Data

AccountID	SalesAmt	...
47	\$1,234.36	...
36	\$2,345.47	...
14	\$3,456.58	...
25	\$4,567.69	...
48	\$5,678.70	...
37	\$6,789.81	...
...	...	...

```
CREATE TABLE ProductSales
WITH (DISTRIBUTION=HASH(AccountID))
AS...
```



# Hash Distributed (AccountID)

## Hash Distributed

ProductSales – Raw Data

AccountID	SalesAmt	...
47	\$1,234.36	...
36	\$2,345.47	...
14	\$3,456.58	...
25	\$4,567.69	...
48	\$5,678.70	...
37	\$6,789.81	...
...	...	...

```
CREATE TABLE ProductSales
WITH (DISTRIBUTION=HASH(AccountID))
AS...
```



# Distributed Data Movement

## Distributed Data Movement – Query

ProductSales

AccountID	SalesAmt	...
47	\$1,234.36	...
36	\$2,345.47	...
14	\$3,456.58	...
25	\$4,567.69	...
48	\$5,678.70	...
37	\$6,789.81	...
...	...	...

SalesAccountTerritory

SATerritoryID	AccountID	...
444	37	...
333	25	...
111	36	...
222	47	...
445	14	...
334	48	...
...	...	...

```
CREATE TABLE ProductSales  
WITH (DISTRIBUTION=HASH(AccountID))  
AS...
```

```
CREATE TABLE SalesAccountTerritory  
WITH (DISTRIBUTION=HASH(SATerritoryID))  
AS...
```

# Distributed Data Movement

## Distributed Data Movement – Query

ProductSales

	AccountID	SalesAmt	...
Node 1:	47	\$1,234.36	...
Node 2:	36	\$2,345.47	...
Node 3:	14	\$3,456.58	...
Node 4:	25	\$4,567.69	...
Node 5:	48	\$5,678.70	...
Node 6:	37	\$6,789.81	...
	...	...	...

SalesAccountTerritory

SATerritoryID	AccountID	...
444	37	...
333	25	...
111	36	...
222	47	...
445	14	...
334	48	...
...	...	...

```
CREATE TABLE ProductSales
WITH (DISTRIBUTION=HASH(AccountID))
AS...
```

```
CREATE TABLE SalesAccountTerritory
WITH (DISTRIBUTION=HASH(SATerritoryID))
AS...
```

# Distributed Data Movement

## Distributed Data Movement – Query

ProductSales

	AccountID	SalesAmt	...
Node 1:	47	\$1,234.36	...
Node 2:	36	\$2,345.47	...
Node 3:	14	\$3,456.58	...
Node 4:	25	\$4,567.69	...
Node 5:	48	\$5,678.70	...
Node 6:	37	\$6,789.81	...
	...	...	...

SalesAccountTerritory

SATerritoryID	AccountID	...
444	37	...
333	25	...
111	36	...
222	47	...
445	14	...
334	48	...
...	...	...

```
CREATE TABLE ProductSales
WITH (DISTRIBUTION=HASH(AccountID))
AS...
```

```
CREATE TABLE SalesAccountTerritory
WITH (DISTRIBUTION=HASH(SATerritoryID))
AS...
```

```
SELECT TOP 25 a.SalesAccountTerritoryName
,TotalSales = SUM(p.SalesAmt)
FROM ProductSales p
JOIN SalesAccountTerritory a
ON a.AccountID = p.AccountID
GROUP BY a.SalesAccountTerritoryName
ORDER BY 2 DESC
```



# Distributed Data Movement

## Distributed Data Movement – Query

ProductSales

	AccountID	SalesAmt	...
Node 1:	47	\$1,234.36	...
Node 2:	36	\$2,345.47	...
Node 3:	14	\$3,456.58	...
Node 4:	25	\$4,567.69	...
Node 5:	48	\$5,678.70	...
Node 6:	37	\$6,789.81	...
	...	...	...

SalesAccountTerritory

SATerritoryID	AccountID	...
444	37	...
333	25	...
111	36	...
222	47	...
445	14	...
334	48	...
...	...	...



AccountID	SATName	...
47	West	...
36	East	...
14	SouthWest	...
25	NorthEast	...
48	South	...

```
CREATE TABLE ProductSales
WITH (DISTRIBUTION=HASH(AccountID))
AS...
```

```
CREATE TABLE SalesAccountTerritory
WITH (DISTRIBUTION=HASH(SATerritoryID))
AS...
```

```
SELECT TOP 25 a.SalesAccountTerritoryName
, TotalSales = SUM(p.SalesAmt)
FROM ProductSales p
JOIN SalesAccountTerritory a
ON a.AccountID = p.AccountID
GROUP BY a.SalesAccountTerritoryName
ORDER BY 2 DESC
```

# Distributed Data Movement

## Distributed Data Movement – Query

ProductSales

	AccountID	SalesAmt	...
Node 1:	47	\$1,234.36	...
Node 2:	36	\$2,345.47	...
Node 3:	14	\$3,456.58	...
Node 4:	25	\$4,567.69	...
Node 5:	48	\$5,678.70	...
Node 6:	37	\$6,789.81	...
	...	...	...

SalesAccountTerritory

	AccountID	SATName	...
Node 1:	47	West	...
Node 2:	36	East	...
Node 3:	14	SouthWest	...
Node 4:	25	NorthEast	...
Node 5:	48	South	...
Node 6:	37	North	...
	...	...	...

```
CREATE TABLE ProductSales
WITH (DISTRIBUTION=HASH(AccountID))
AS...

CREATE TABLE SalesAccountTerritory
WITH (DISTRIBUTION=HASH(SATerritoryID))
AS...

SELECT TOP 25 a.SalesAccountTerritoryName
,TotalSales = SUM(p.SalesAmt)
FROM ProductSales p
JOIN SalesAccountTerritory a
ON a.AccountID = p.AccountID
GROUP BY a.SalesAccountTerritoryName
ORDER BY 2 DESC
```



# Distributed Data Movement

## Distributed Data Movement – Query

ProductSales

SalesAccountTerritory

	AccountID	SalesAmt	...	AccountID	SATName	...
Node 1:	47	\$1,234.36	...	47	West	...
Node 2:	36	\$2,345.47	...	36	East	...
Node 3:	14	\$3,456.58	...	14	SouthWest	...
Node 4:	25	\$4,567.69	...	25	NorthEast	...
Node 5:	48	\$5,678.70	...	48	South	...
Node 6:	37	\$6,789.81	...	37	North	...
	...	...	...	...	...	...

SATName	TotalSales
North	\$6,789.81
South	\$5,678.70
NorthEast	\$4,567.69
SouthWest	\$3,456.58
East	\$2,345.47
West	\$1,234.36

```
CREATE TABLE ProductSales
WITH (DISTRIBUTION=HASH(AccountID))
AS...
```

```
CREATE TABLE SalesAccountTerritory
WITH (DISTRIBUTION=HASH(SATerritoryID))
AS...
```

```
SELECT TOP 25 a.SalesAccountTerritoryName
              ,TotalSales = SUM(p.SalesAmt)
FROM ProductSales p
JOIN SalesAccountTerritory a
ON a.AccountID = p.AccountID
GROUP BY a.SalesAccountTerritoryName
ORDER BY 2 DESC
```

# Distribution Guidance

For large fact tables, best option is to Hash Distribute


- Clustered Columnstore
- Distribute on column that is joined to other fact tables or large dimensions
- Primary or surrogate key maybe a good choice for distribution

However, be mindful of ...


- Hash column should have highly distinct values (Minimum 600 distinct values)
- Avoid distributing on a date column
- Avoid distributing on column with high frequency of NULLs and default values (e.g. -1)
- Distribution column is NOT updatable
- For compatible joins use the same data types for two distributed tables

If there are no distribution columns that make sense, then use Round Robin as last resort

# Skewed Distribution



# Replicated Tables



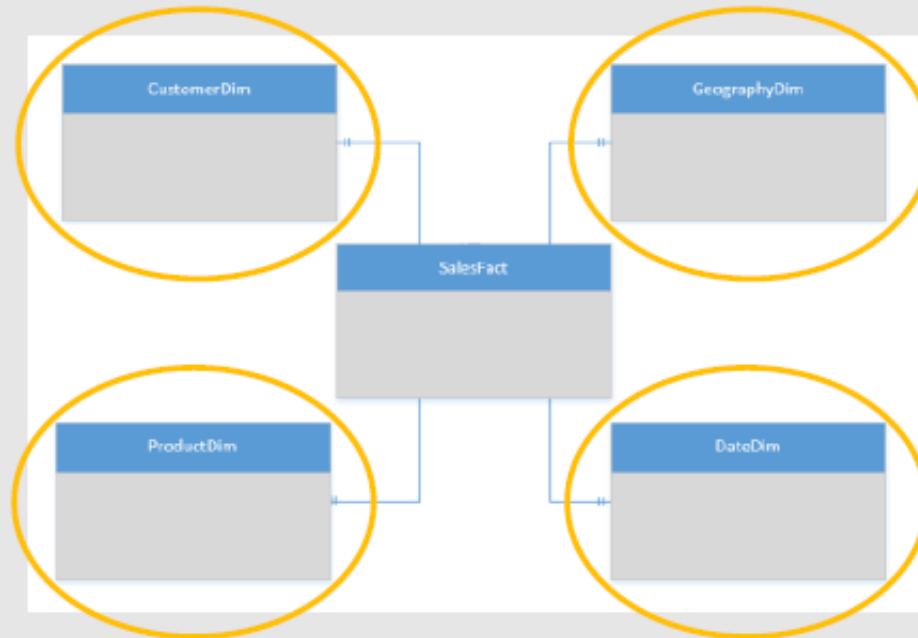
# Replicated Tables Scenarios

## Scenarios to consider using Replicated tables:

Star schema reporting


Dimensions – descriptive entities about fact data

ETL master data, common domain data used during transaction loading



# Replicated Tables .. How it works


## Replicated Tables... How it works





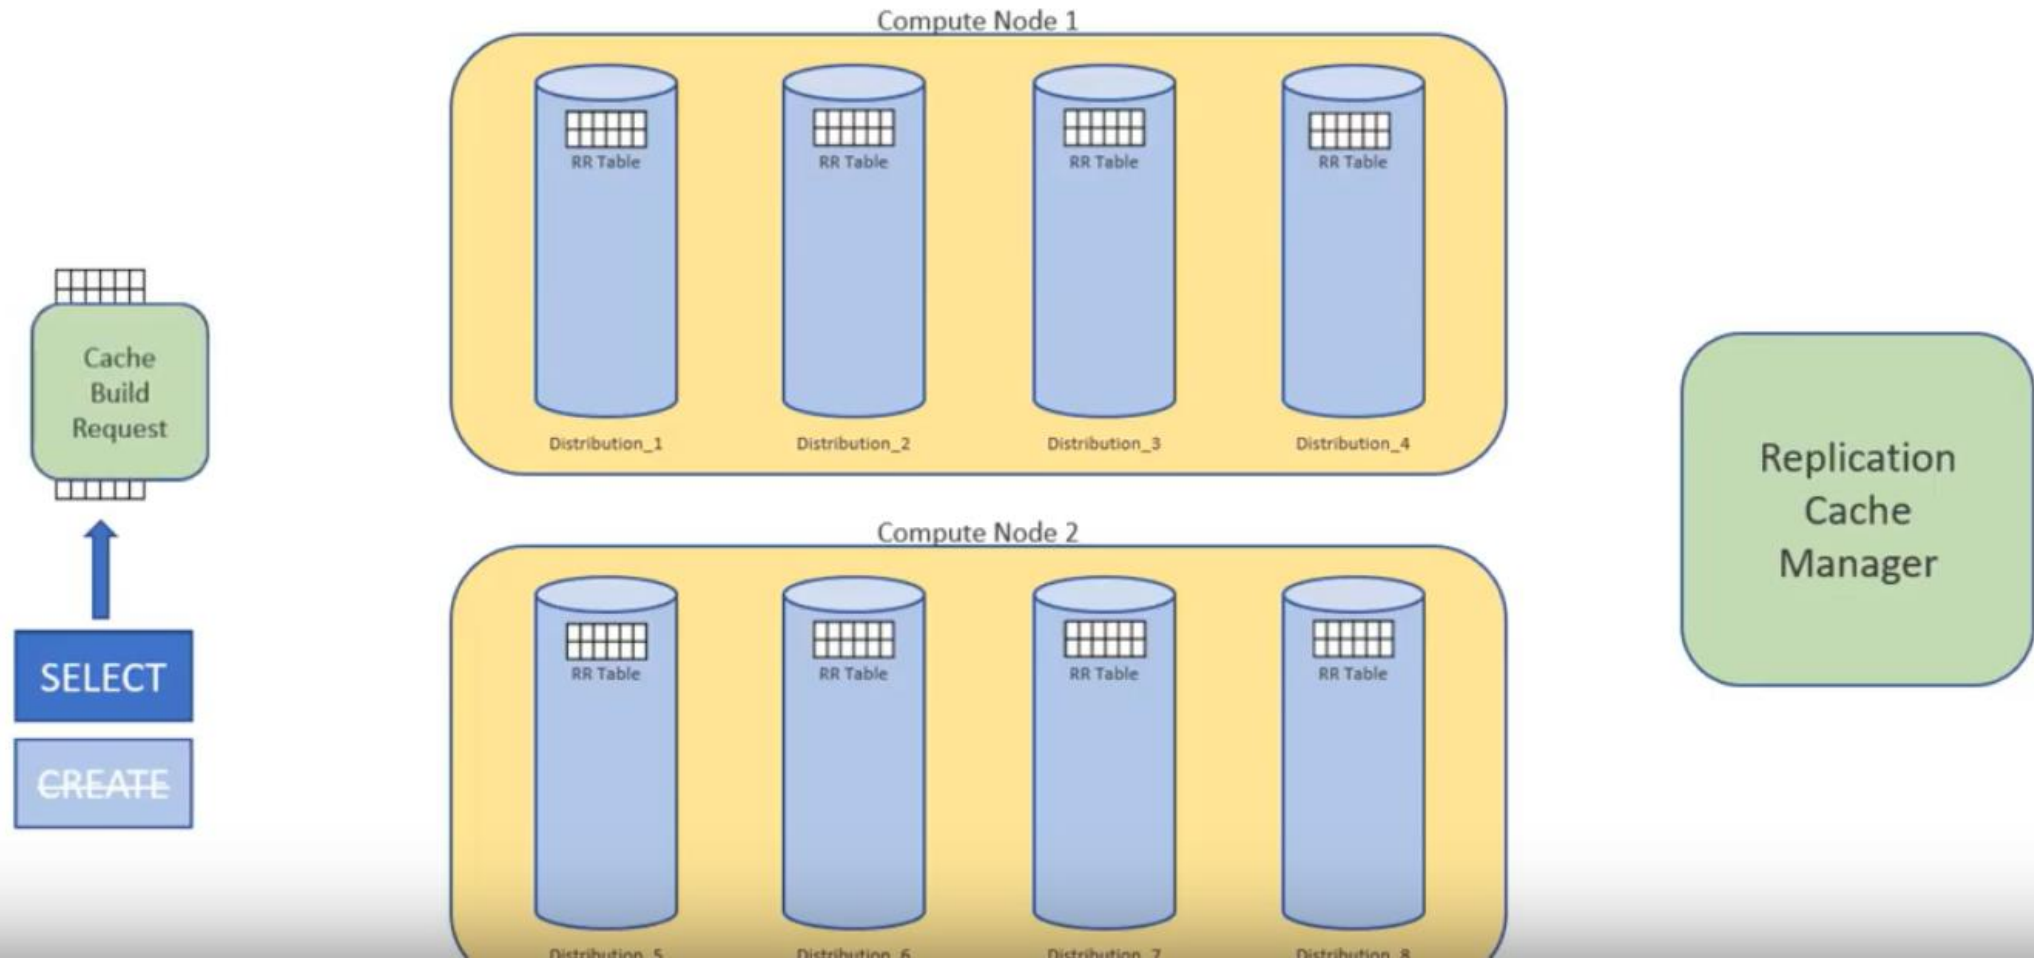
# Replicated Tables .. How it works

## Replicated Tables... How it works



# Replicated Tables .. How it works

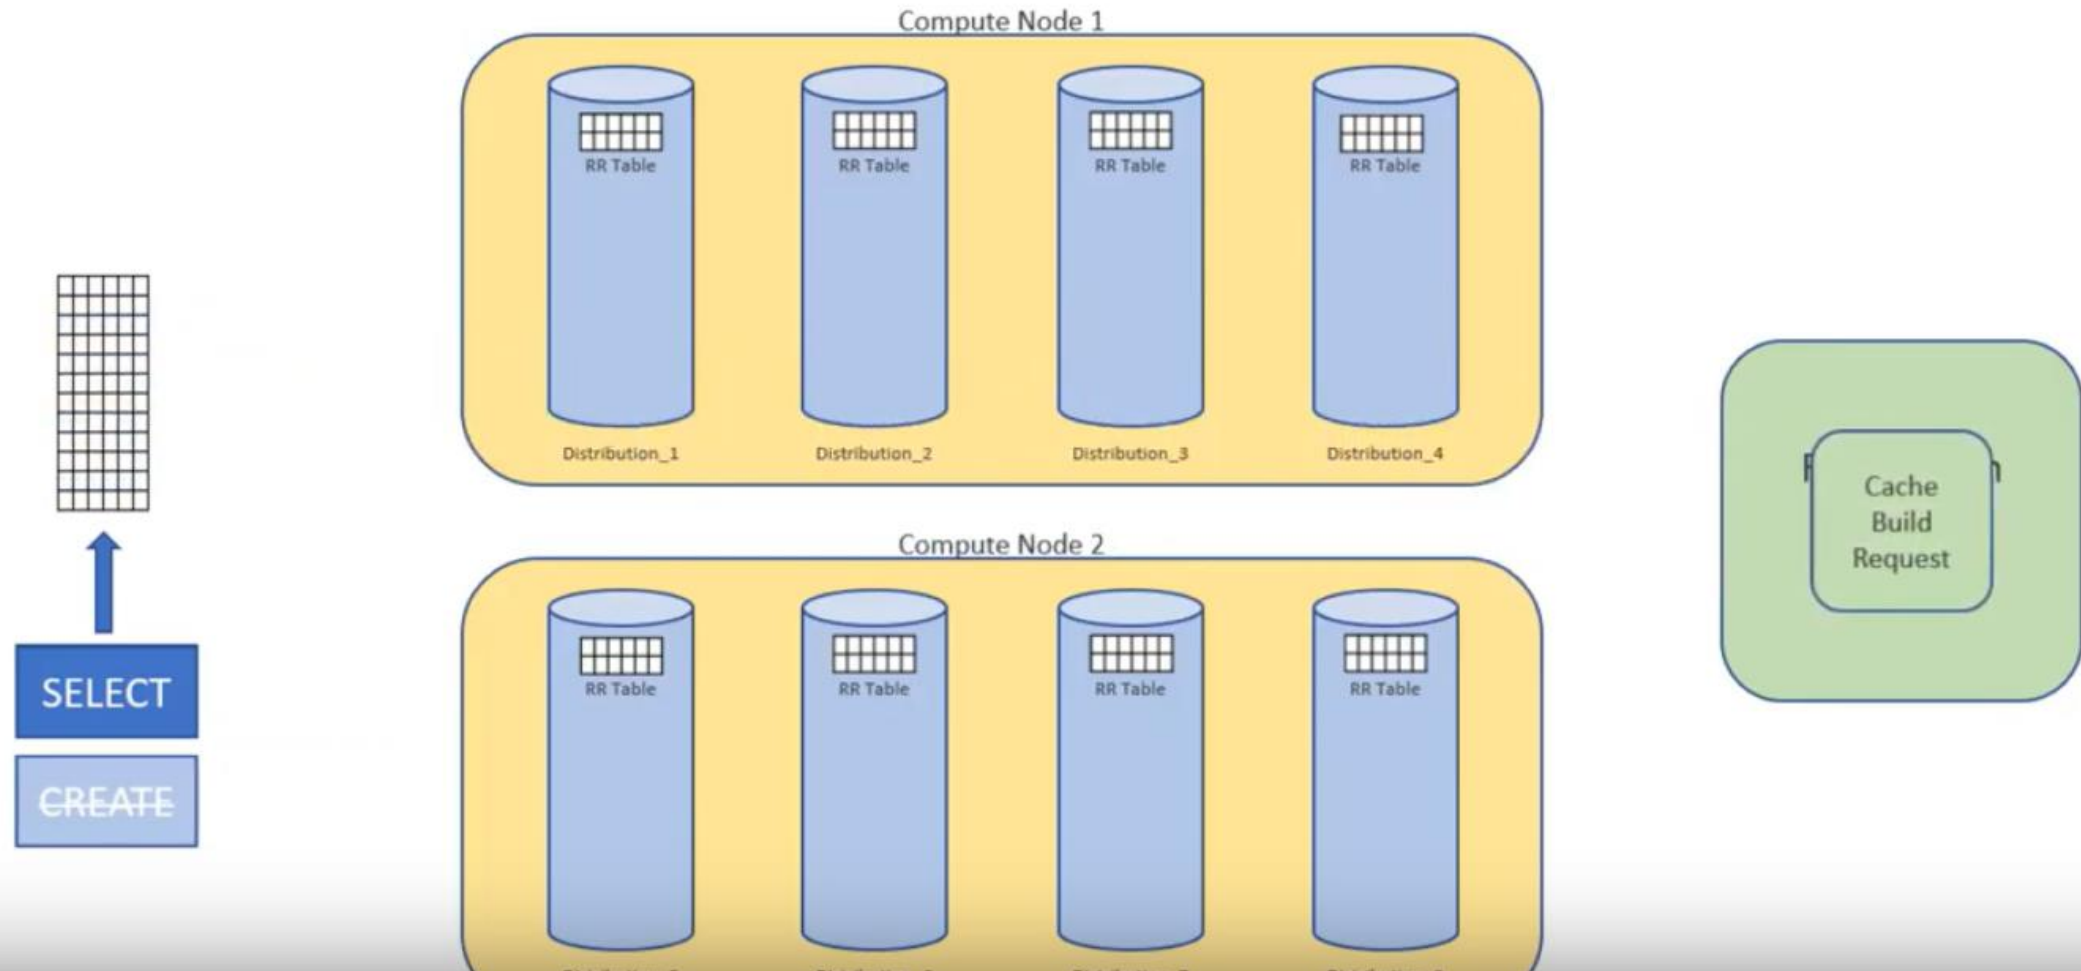
## Replicated Tables... How it works






# Replicated Tables .. How it works

## Replicated Tables... How it works

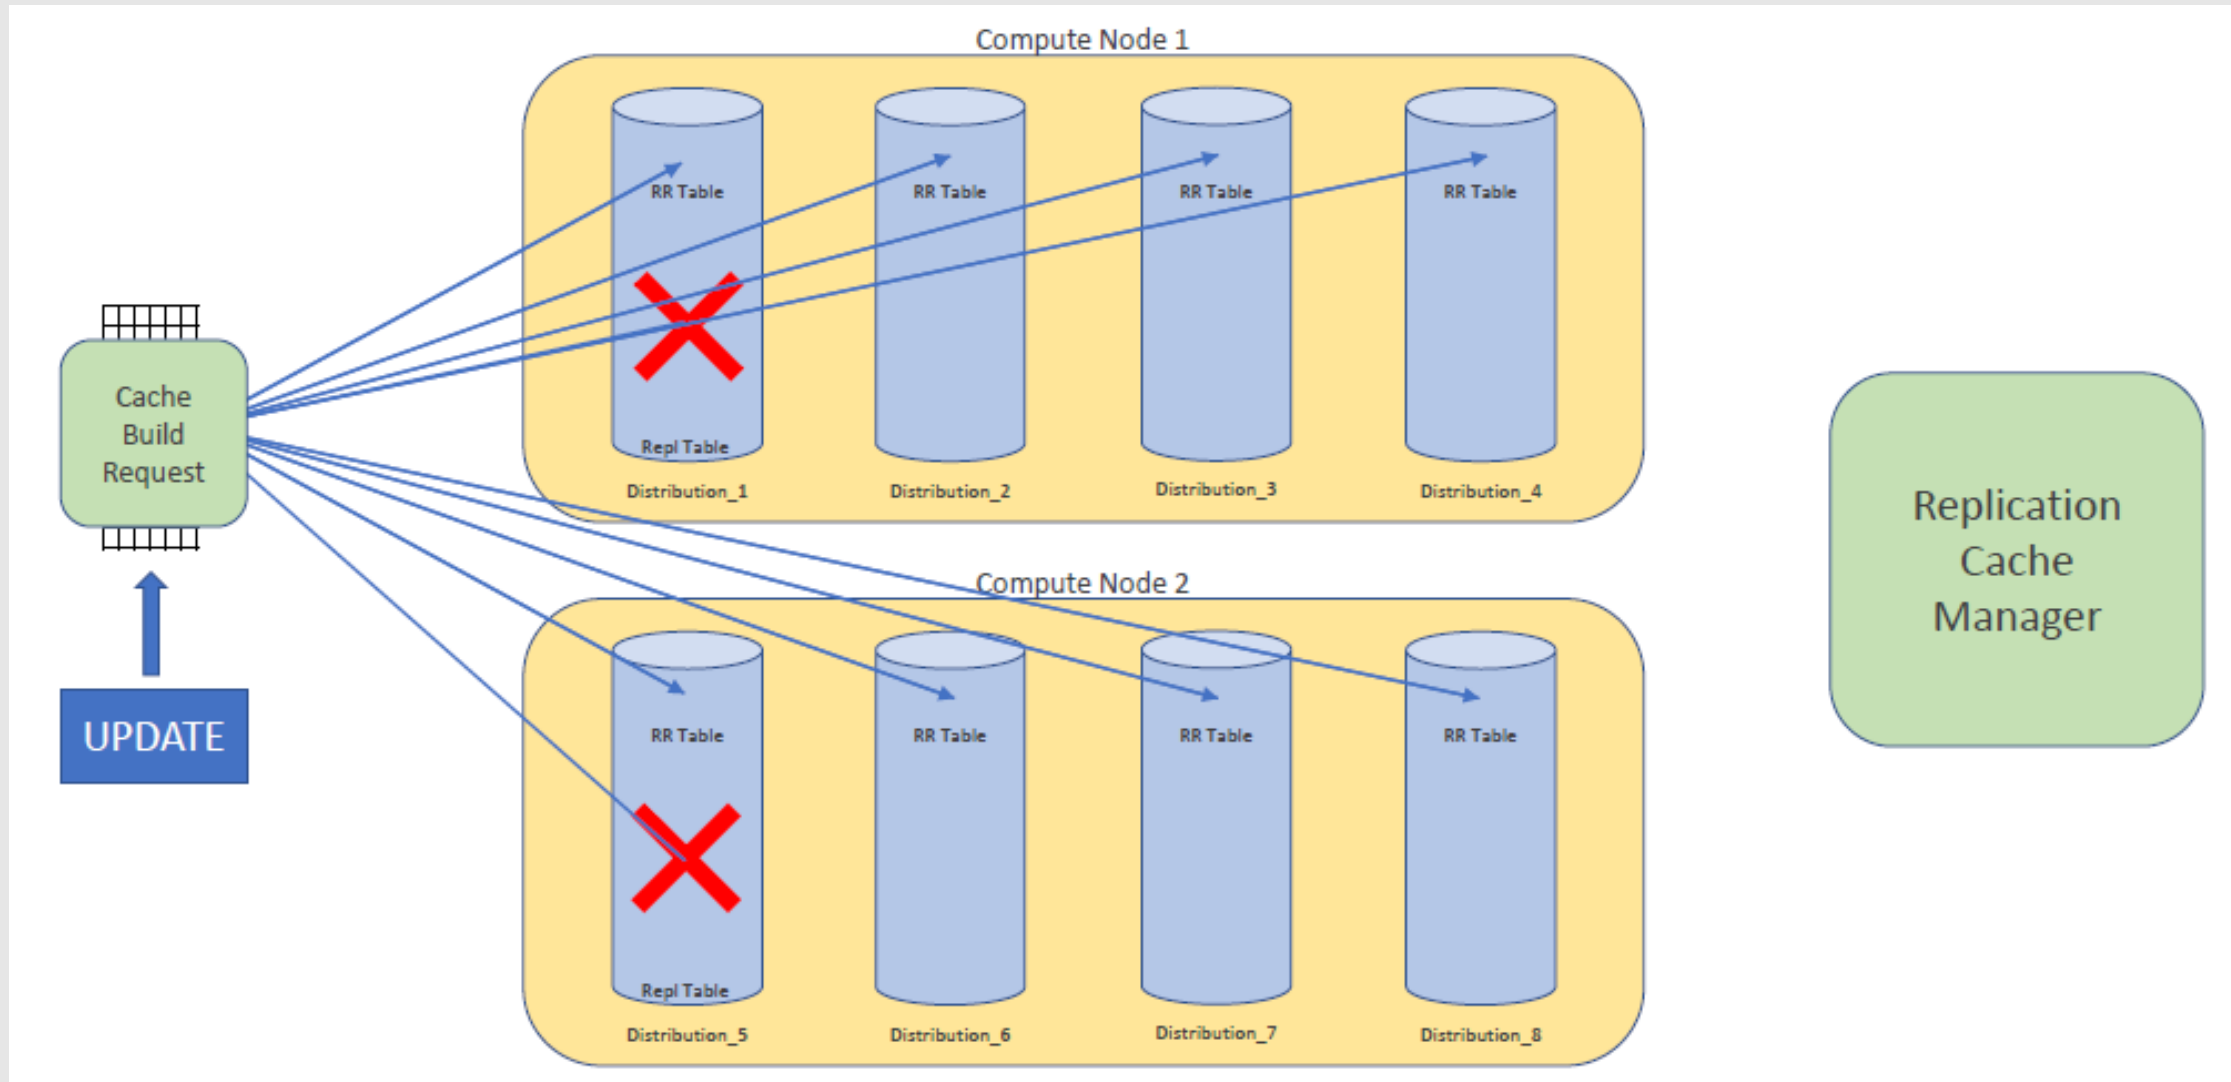


# Replicated Tables .. How it works

## Replicated Tables... How it works




# Replicated Tables .. How it works



# Replicated Tables Best Practices

- Good for
  - Tables used frequently in Joins
  - Tables size less than 2GB on disk
  - Queries with simple predicates (Ex: Equal, not equal)
- Not good for
  - Tables with frequent modifications
  - Frequent scaling of DW
  - Large Tables (>2GB)
  - Tables with large number of columns (but query small number of columns)

# Concurrency, Concurrency Slots and Resource Classes

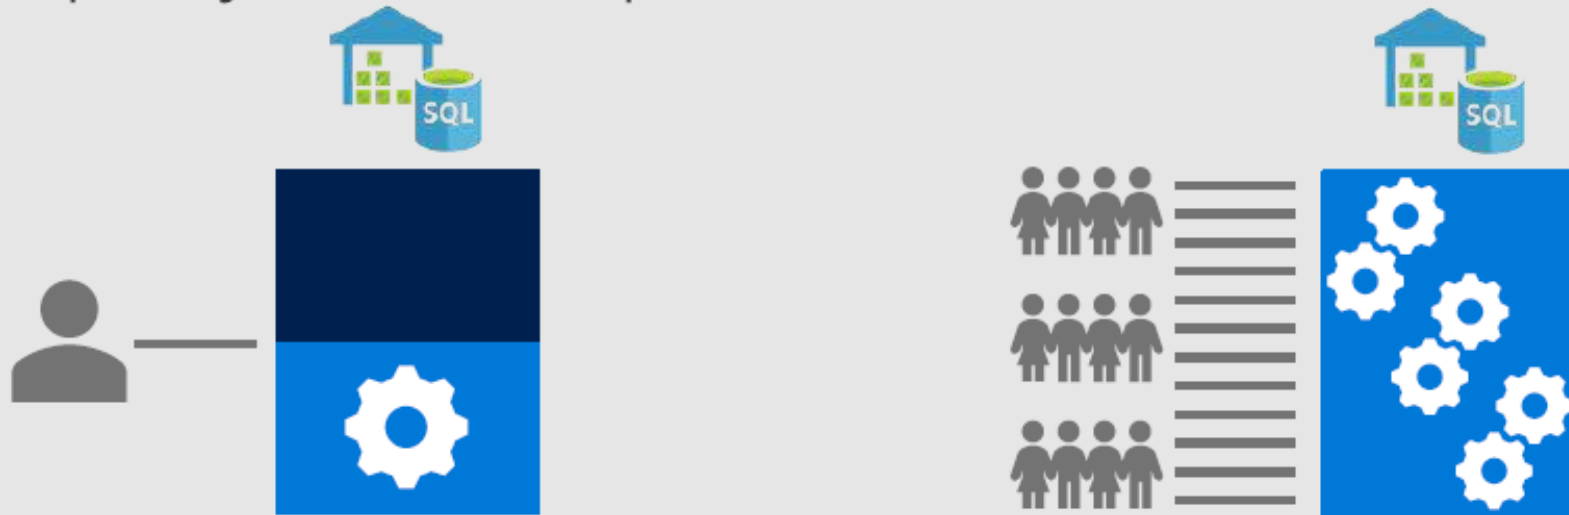


# Importance of Concurrency

Parallel queries = more throughput (even on large, scale-out DWs)

Mixed workloads (transform, load, export, query)

A single query is not expected to consume all resources



# Concurrency Slots

## SQL DW concurrency concepts

Service Level

DW6000c

Concurrency slots

Memory allocated to resource classes


Resource Class

smallrc

staticrc10

# Concurrency Slots


Concurrency slots






# Concurrency Slots

## Concurrency slots




Total system  
available memory



Concurrency Slots  
20 per DW500c  
4 per DW100

# Concurrency Slots

## Concurrent Query and Slots



# Resource Classes

## Resource Classes

Resource classes map concurrency slots to user requests and thus, allow control over **memory** allocation.

There are two types of resource classes: **Static** and **Dynamic**

# Resource Classes

## Resource Classes – Dynamic

Allocates variable amounts of memory depending on the scale of the DW instance.

- ✓ Beneficial for variable sized workloads that scale to meet demand.
- ✗ There is no increase in concurrency with scaling. Should be avoided.



Scaling up →

# Resource Classes

## Resource Classes – Static

Allocates a fixed amount of memory regardless of the scale level.

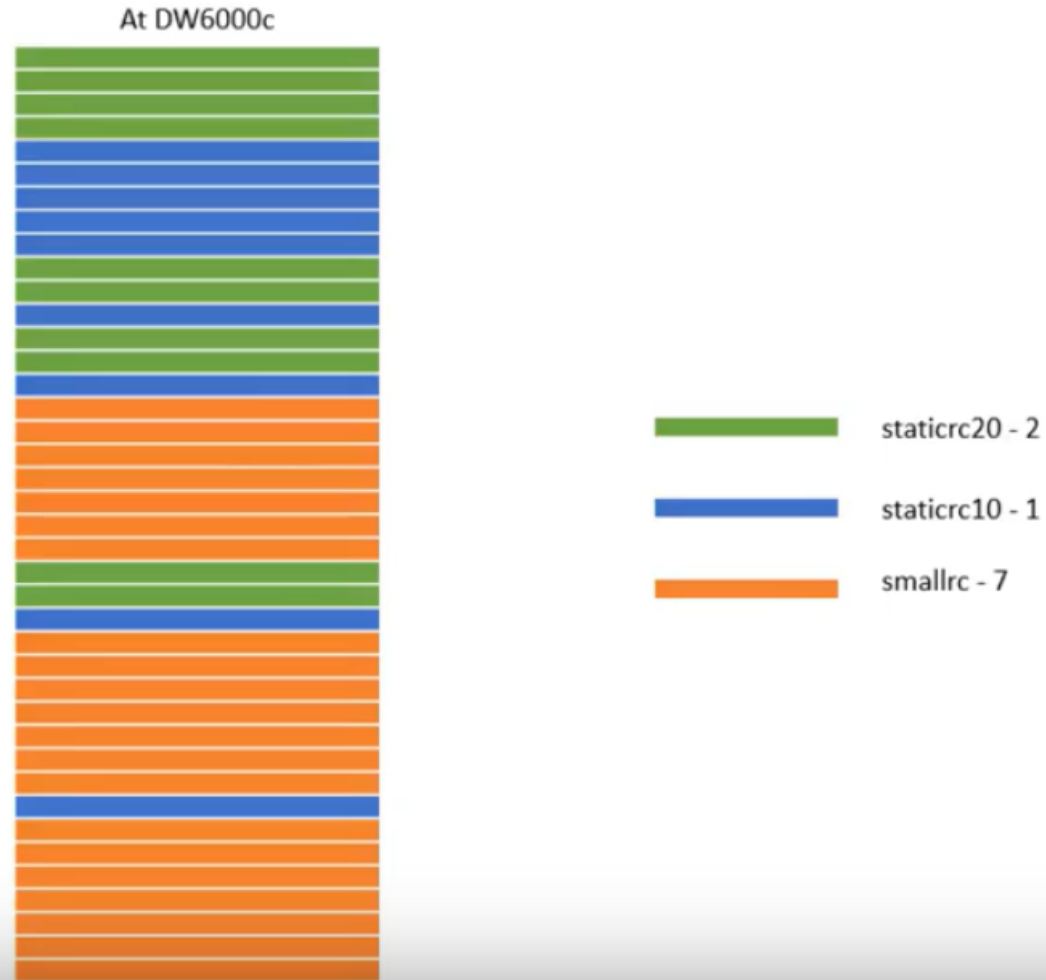
- ✓ Essential for high query concurrency workloads.
- ✗ Queries may run the same regardless of the service level.



Scaling up →

# Resource Classes


## Bringing it all together



# Resource Class Best Practices

- Be aware of the user binding
- Use right resource class for right workload
- Prefer Static for higher concurrency
- Prefer Dynamic for increased perf with scale
- Be aware of the RC precedence
  - Dynamic over Static
  - Higher RC over lower
- Explicitly remove user from unnecessary RCs
- Use custom coding for automating RC assignment


Index/Stats





# Clustered Columnstore Indexes(CCI) Best Practices

- Be aware of the default
- Not efficient for
  - transient data (frequent updates/deletes)
  - Singleton loads or micro-batches
  - Small tables (<100 Million rows)
- >100K rows per rowgroup
- Highest possible RC for loading
  - Calculate memory requirements
    - $72MB + (\#rows * \#columns * 8B) + (\#rows * \#SSC * 32B) + (\#LSC * 16MB)$
- Reduce memory requirements
  - Small number of columns/table
  - Reduce columns with string data type
  - Do not over partition
  - Simplify load queries
  - Adjust MAXDOP (if needed)



# Index/Stats Best Practices

- Perform regular CCI health checks to monitor
  - # of open row groups
  - # of rows per row group (100 K to 1 Million)
  - Reason for trimming
- Perform regular Index Maintenance
  - Rebuild/Reorganize
  - Partition Rebuild
  - CTAS/Partition switch if Needed
- Statistics
  - Use auto-create stats option
  - Create multi-column stats as needed
  - Update stats immediately after large data modifications
  - Auto-update stats option (coming soon)

# Monitoring



# Monitoring Options

- Azure Monitor
  - Insights
  - Alerts
  - Dashboard
  - Views
  - PBI Integration
- Query Store
- Operations Management Suite (OMS) integration

## Command-line

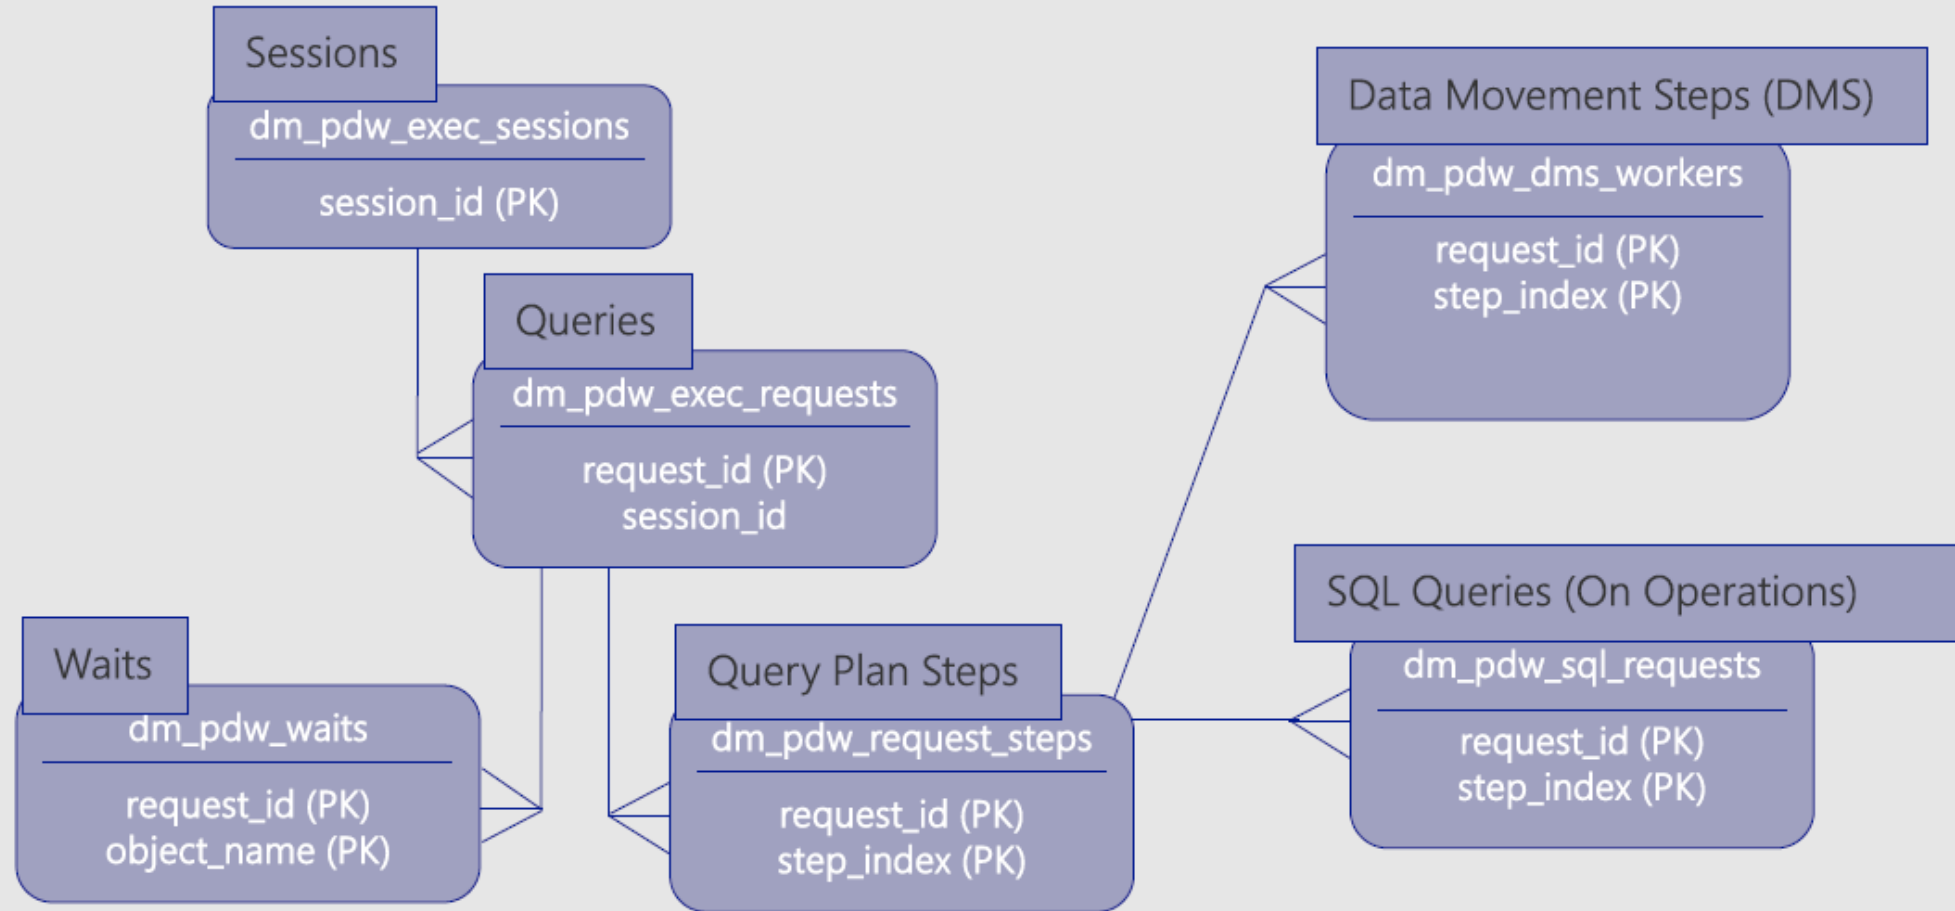
- Rich set of DMVs
- Lower level of diagnostics

# Azure Portal

## Key Metrics

- CPU Percentage
- IO Percentage
- DWU (limit, used and percentage used)
- Connections (successful, failed, blocked by firewall)
- Gen2: Cache metrics (used and cache hit percentage)

# Execution DMVs



# Additional Resources

SQL Data Warehouse Documentation - <https://docs.microsoft.com/en-us/azure/sql-data-warehouse/>



# Q&A

