



Azure Data Warehouse In-A-Day

Keshav Pokkuluri
Casey Karst
Kal Yella



Agenda

8:45 AM to 9:00 AM	Welcome
9:00 AM to 9:45 AM	Datawarehouse Patterns & SQL DW Overview
9:45AM to 10:30 AM	SQL DW Gen2 New Features & Functionality

Break: 10:30 AM to 10:45 AM

10:45 AM to 11:15 PM	SQLDW Loading Best Practices
11:15 AM to 12:00 PM	SQLDW Operational Best Practices

Lunch: 12:00 PM to 1:00 PM SQLDW Roadmap Discussion

1:00 PM to 1:45 PM	SQL DW Compete Discussion
--------------------	---------------------------


Break: 1:45 PM to 2:00 PM

2:00 PM to 2:45 PM	Lab 1: Data loading scenarios and best practices
2:45 PM to 3:30 PM	Lab 2: Performance Tuning best practices
3:30 PM to 4:30 PM	Lab 3: Monitoring, Maintenance and Security
4:30 PM to 5:00 PM	Q&A and Wrap-up

Azure SQL DW Service

A relational data warehouse-as-a-service, fully managed by Microsoft.
Industries first elastic cloud data warehouse with proven SQL Server capabilities.
Support your smallest to your largest data storage needs.

Elastic scale & performance




- Scales beyond petabytes of data
- Cloud DW Scale-out Processing
- Instant-on compute scales in seconds
- Query Relational / Non-Relational

Powered by the Cloud

Get started in minutes

Integrated with Azure ML, PowerBI, ADB & ADF



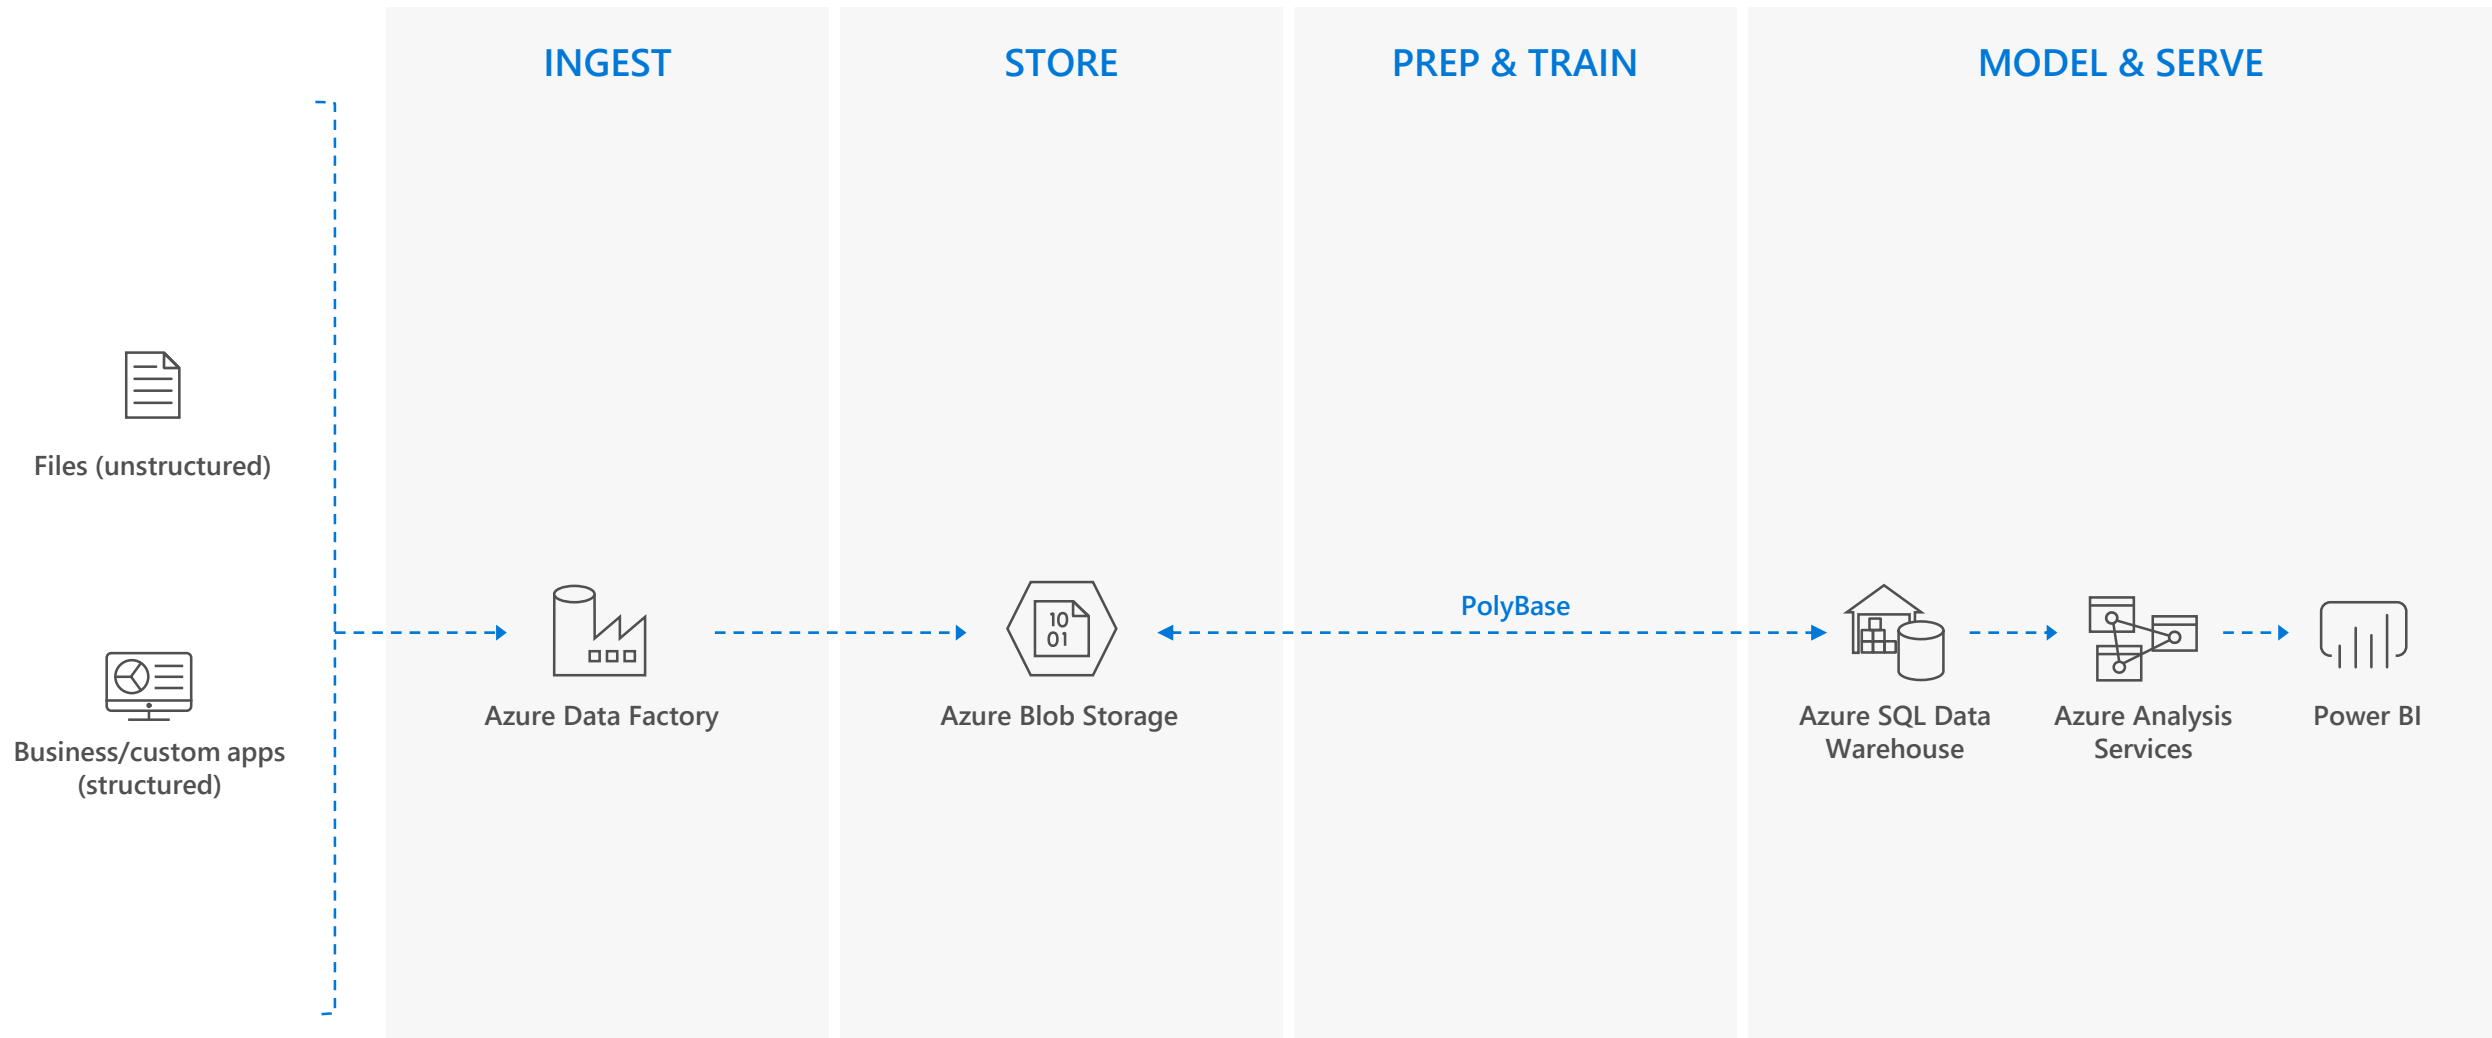
Market Leading Price & Performance



Simple billing compute & storage

Pay for what you need, when you need it
with dynamic pause

CLOUD DATA WAREHOUSE



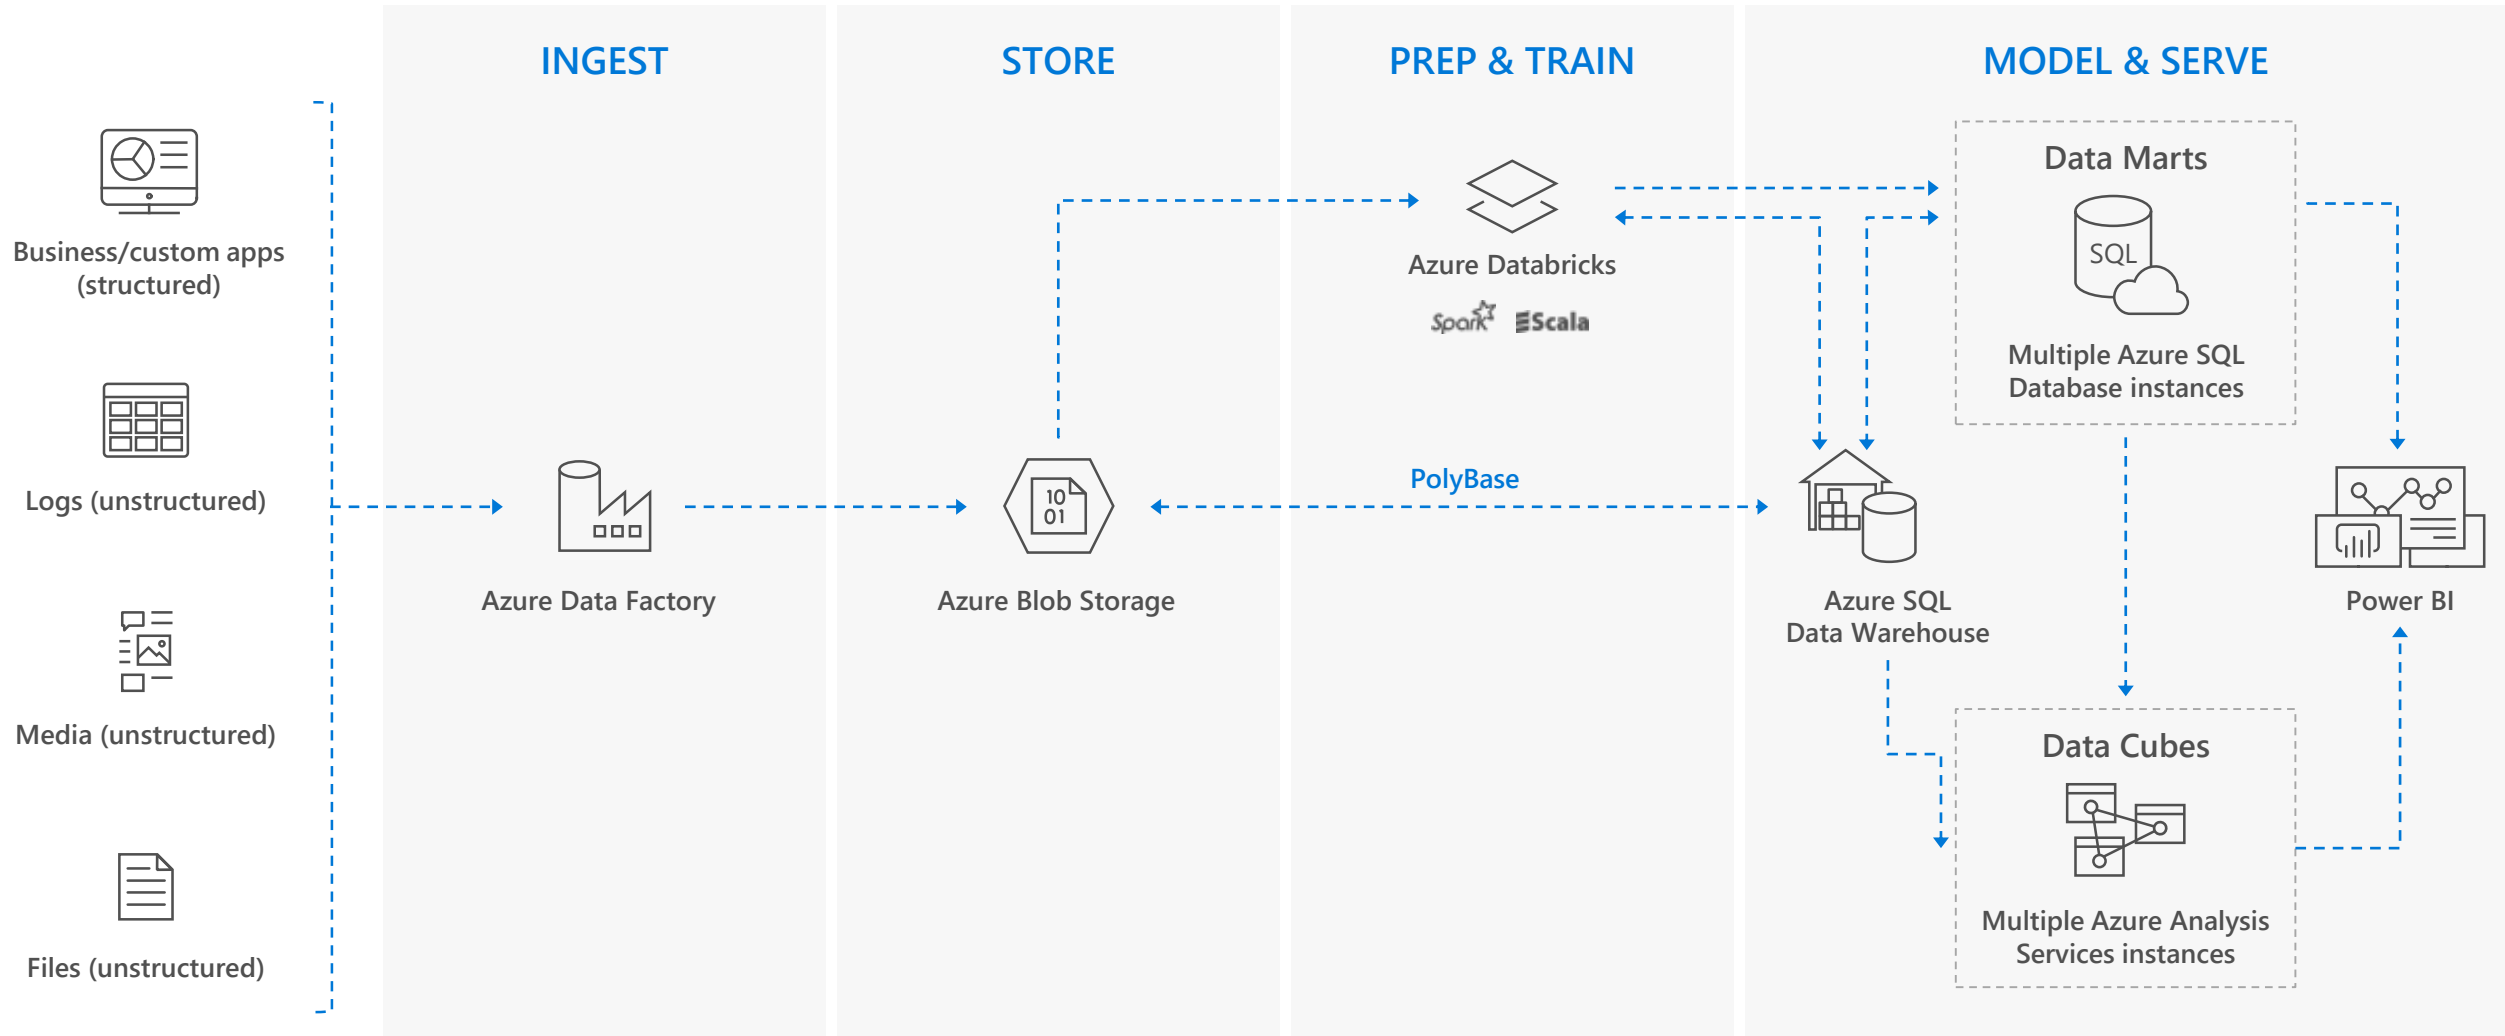
Microsoft Azure also supports other Big Data services like Azure HDInsight and Azure Data Lake to allow customers to tailor the above architecture to meet their unique needs.

MODERN DATA WAREHOUSE



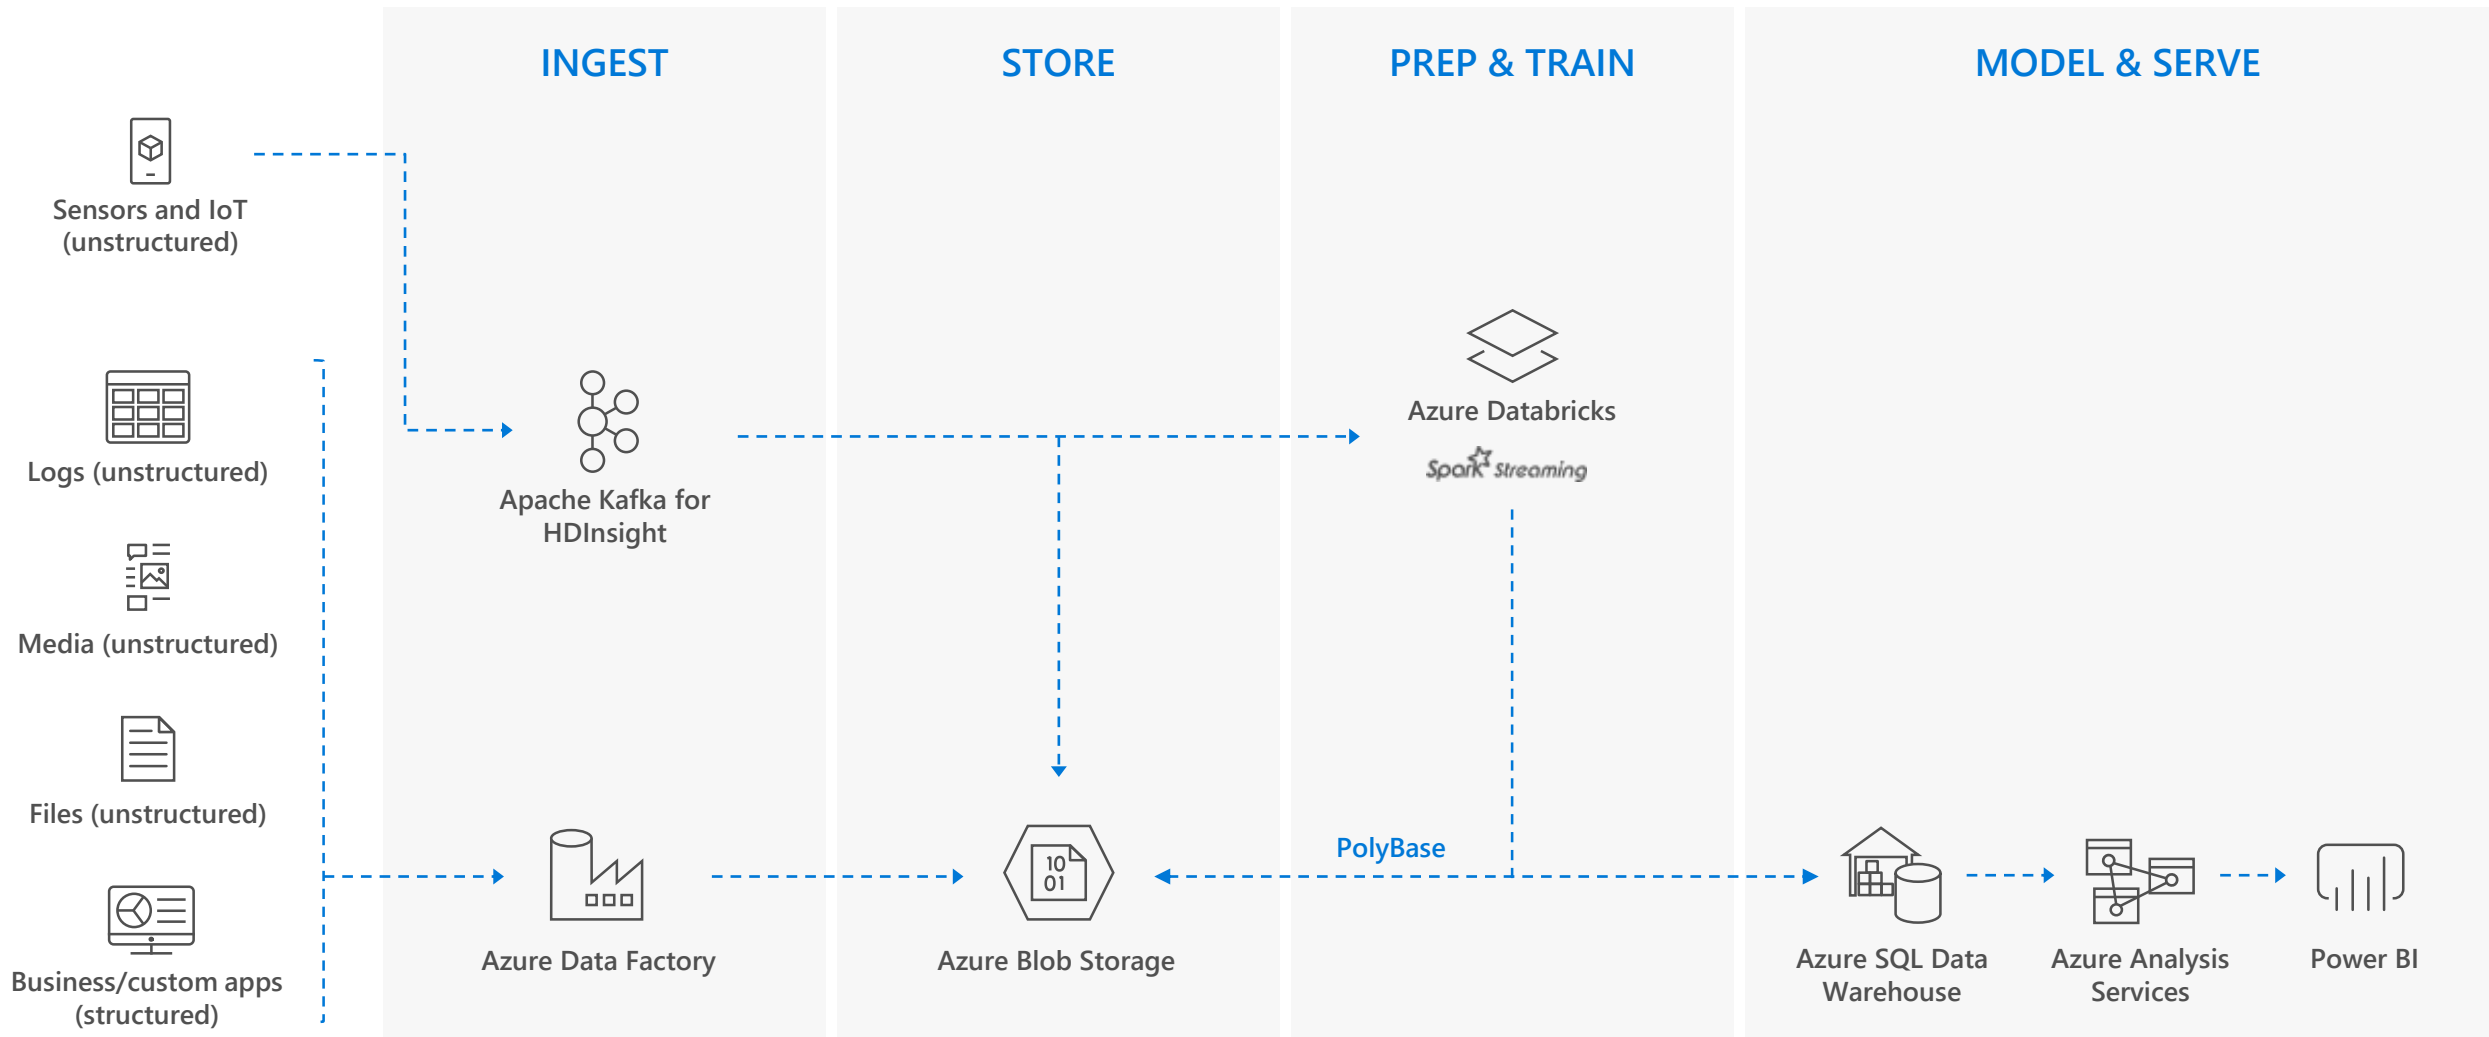
Microsoft Azure also supports other Big Data services like Azure HDInsight and Azure Data Lake to allow customers to tailor the above architecture to meet their unique needs.

HUB & SPOKE ARCHITECTURE FOR BI



Microsoft Azure supports other services like Azure HDInsight and Azure Data Lake in various layers to allow customers a truly customized solution.

REAL TIME ANALYTICS




Microsoft Azure also supports other Big Data services like Azure IoT Hub, Azure Event Hubs, Azure Machine Learning and Azure Data Lake to allow customers to tailor the above architecture to meet their unique needs.

Azure SQL DW Overview


Azure SQL DW CSE



Under the hood - Architecture



SQL DW Fundamentals



Control
Connection and tool endpoint.
Coordinates compute activity.


Compute
Handles query processing,
ability to scale up/down

DMS: Data Movement Services
Coordinates data movement
between nodes

Storage
Remote storage. Scales
independently of compute

Data Warehouse Units

Normalized amount of compute
Converts to billing units i.e. what you pay



DWUc	Nodes	Dist/ Node
100	1	60
200	1	60
300	1	60
400	1	60
500	1	60
1000	2	30
1500	3	20
2000	4	15
2500	5	12
3000	6	10
5000	10	6
...		
30000	60	1

Separate compute from storage



Compute

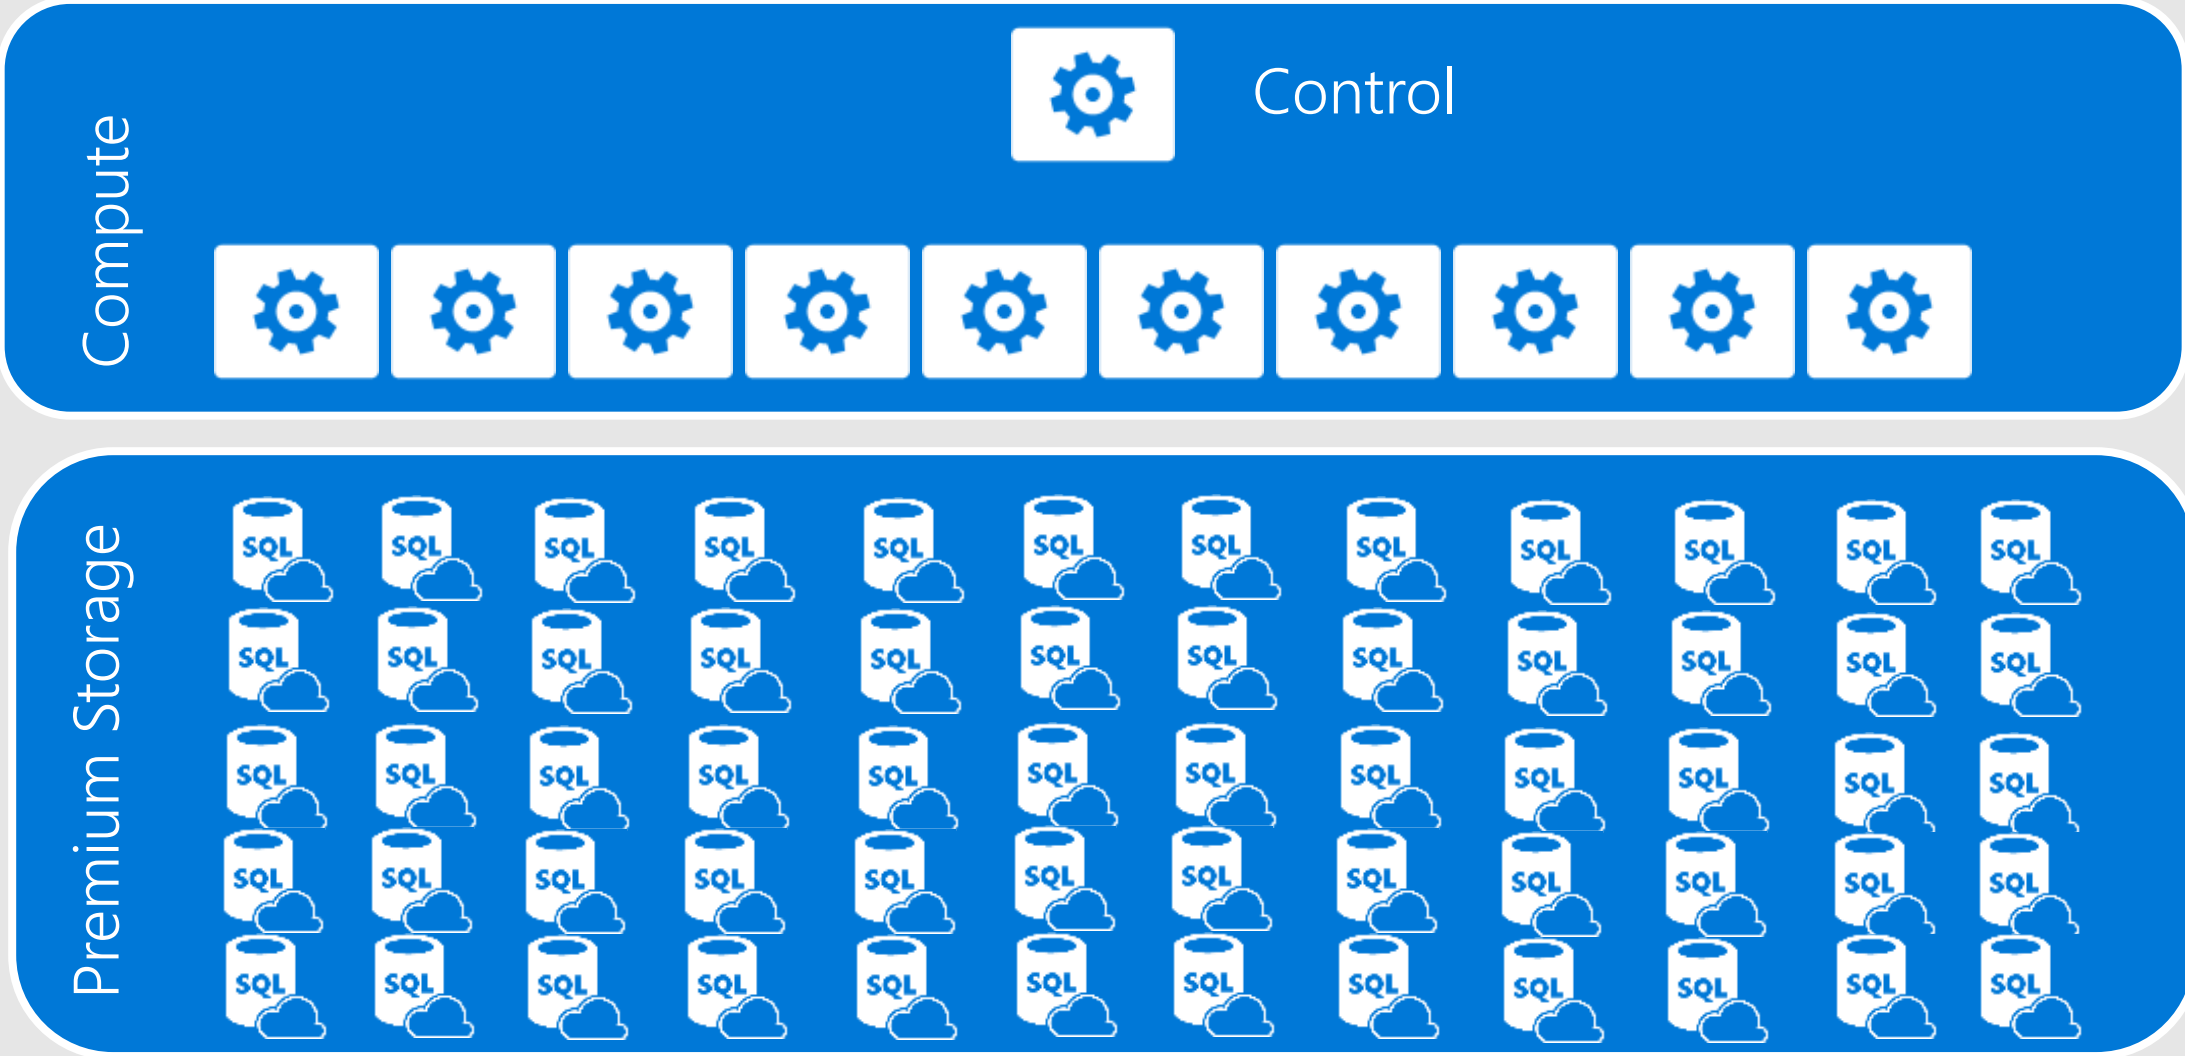


Control

Premium Storage



Independently scale compute



Pause and resume workload



Compute



Control


Premium Storage




Query Execution



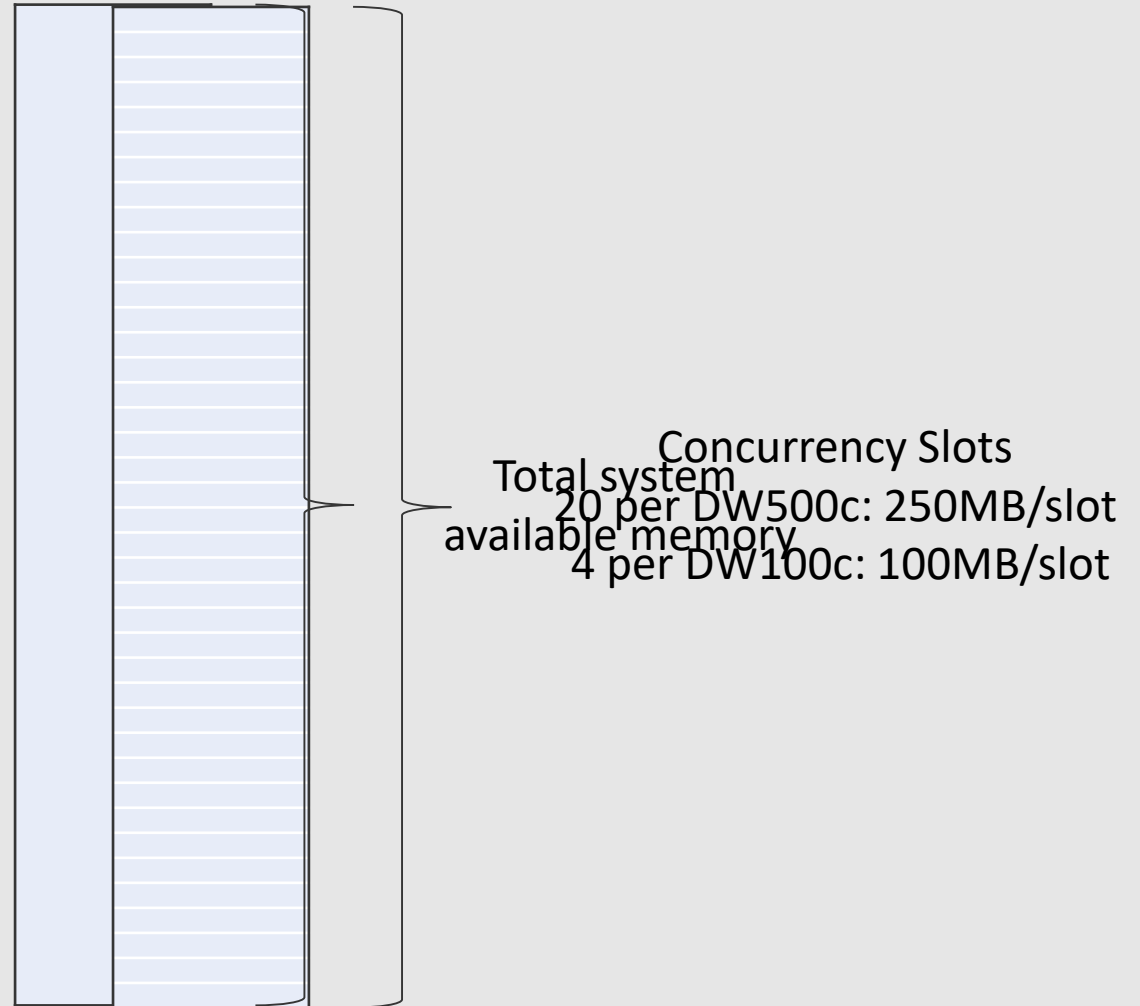
Result



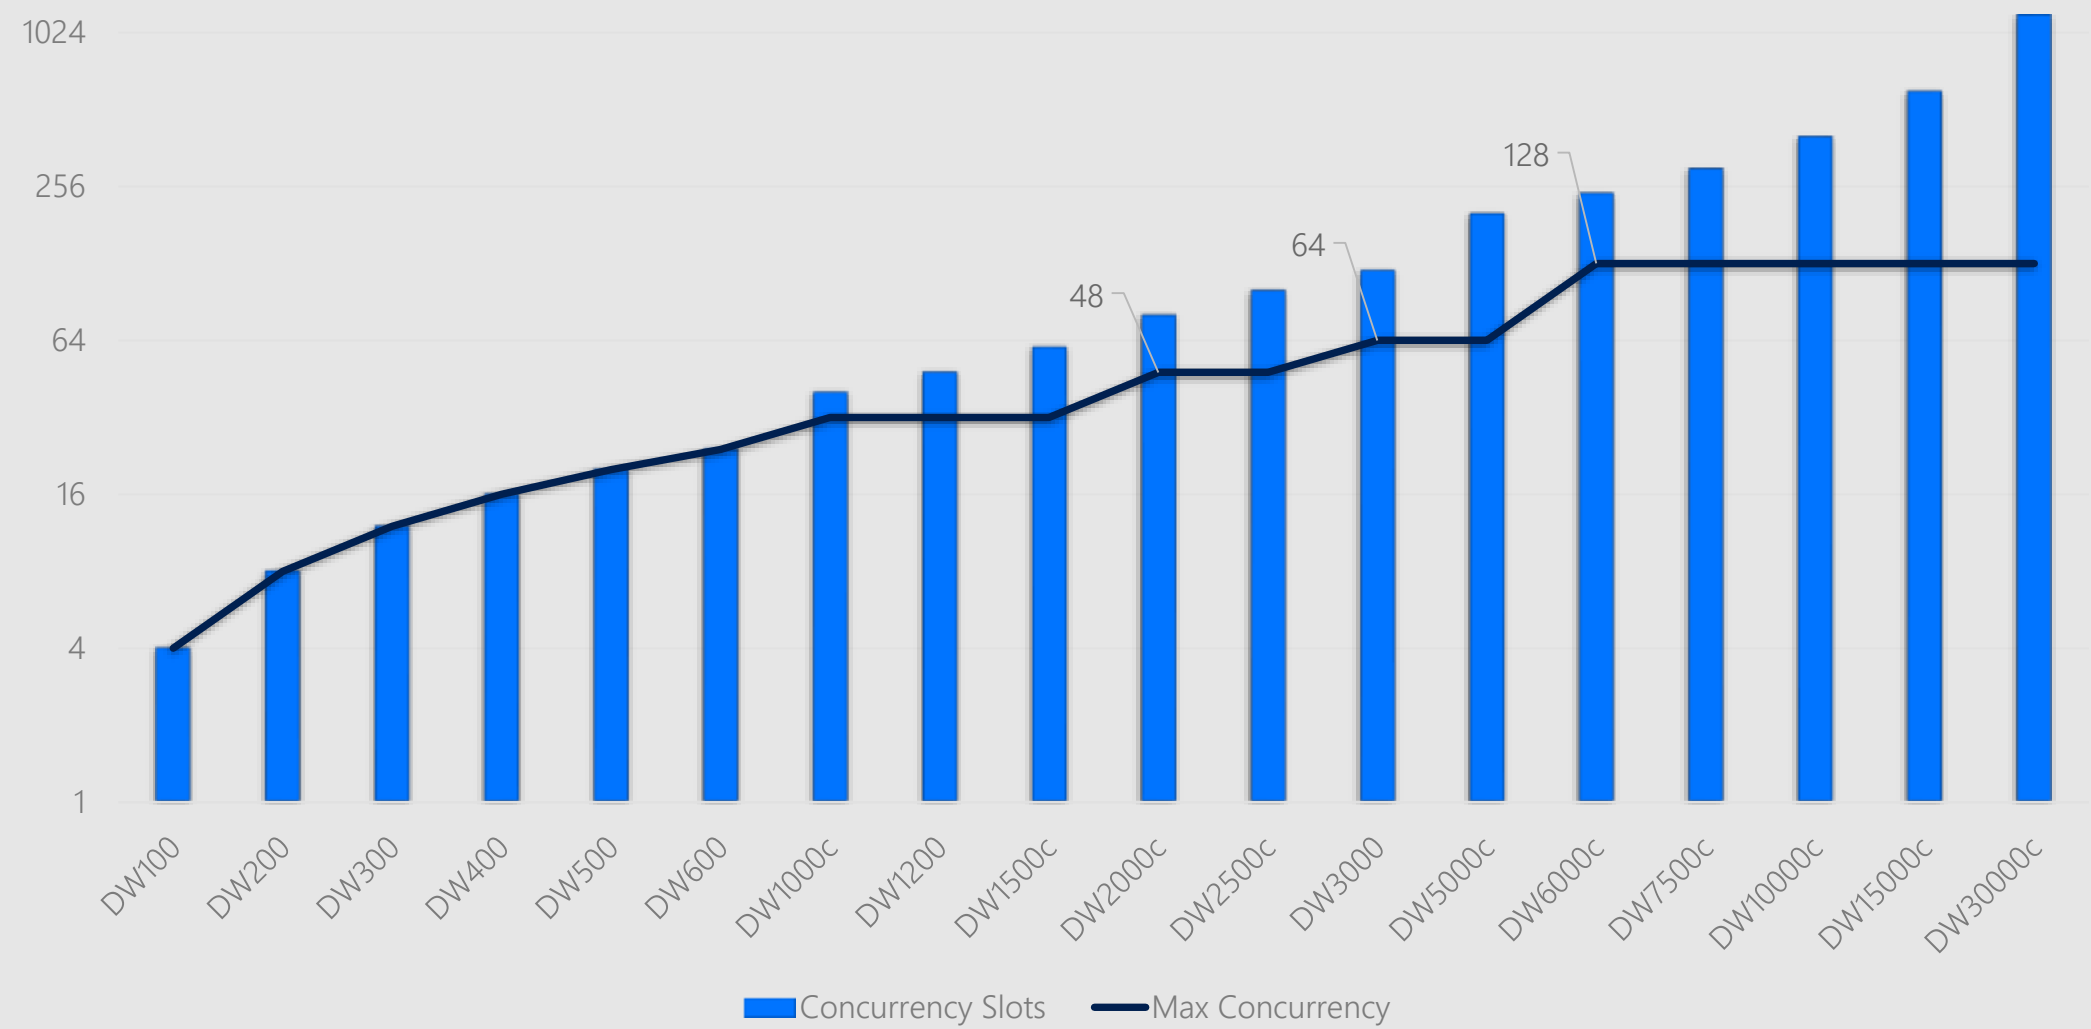
Concurrent queries



Concurrency slots



Concurrent Query and Slots



Resource classes

Dynamic

Increases resource consumption as you scale

No increase in concurrency as you scale

Static

Maintain resource consumption as you scale

Increase concurrent queries as you scale


Consume Slots

Increase memory
Isolate resources

Resource Classes – Dynamic

Allocates variable amounts of memory depending on the scale of the DW instance.

- ✓ Beneficial for variable sized workloads that scale to meet demand.
There is no increase in concurrency with scaling.



Scaling up →

Dynamic Resource Classes

With Gen2, dynamic resource pools were introduced with a 3-10-22-70 model for resource allocations.

Resource Class	Percent Resources	Concurrency
SmallRc	3%	32
MediumRc	10%	10
LargeRc	22%	4
XLargeRc	70%	1

Resource Classes – Static

Allocates a fixed amount of memory regardless of the scale level.

- ✓ Essential for high query concurrency workloads.
Queries may run the same regardless of the service level.



Scaling up →

Gen 2 Concurrency – Static RC

Service Level	Maximum concurrent queries	Concurrency slots available	staticrc10	staticrc20	staticrc30	staticrc40	staticrc50	staticrc60	staticrc70	staticrc80
DW100c	4	4	1	2	4	4	4	4	4	4
DW200c	8	8	1	2	4	8	8	8	8	8
DW300c	12	12	1	2	4	8	8	8	8	8
DW400c	16	16	1	2	4	8	16	16	16	16
DW500c	20	20	1	2	4	8	16	16	16	16
DW1000c	32	40	1	2	4	8	16	32	32	32
DW1500c	32	60	1	2	4	8	16	32	32	32
DW2000c	48	80	1	2	4	8	16	32	64	64
DW2500c	48	100	1	2	4	8	16	32	64	64
DW3000c	64	120	1	2	4	8	16	32	64	64
DW5000c	64	200	1	2	4	8	16	32	64	128
DW6000c	128	240	1	2	4	8	16	32	64	128
DW7500c	128	300	1	2	4	8	16	32	64	128
DW10000c	128	400	1	2	4	8	16	32	64	128
DW15000c	128	600	1	2	4	8	16	32	64	128
DW30000c	128	1200	1	2	4	8	16	32	64	128

Data Distribution




Table Distribution Options

Hash Distributed

Data divided across nodes based on hashing algorithm

Same value will always hash to same distribution

Single column only

Check for Data Skew,
NULLS, -1

Round Robin (Default)

Data distributed evenly across nodes

Easy place to start, don't need to know anything about the data

Simplicity at a cost

Will incur more data movement at query time

Replicated

Data repeated on every node

Simplifies many query plans and reduces data movement

Best with joining hash table

Consumes more space
Joining two Replicated
Table runs
on one node

Creating tables

```
CREATE TABLE [build].[FactOnlineSales]
```

```
(  
    [OnlineSalesKey]      int      NOT NULL  
, [DateKey]              datetime NOT NULL  
, [StoreKey]             int      NOT NULL  
, [ProductKey]           int      NOT NULL  
, [PromotionKey]         int      NOT NULL  
, [CurrencyKey]          int      NOT NULL  
, [CustomerKey]          int      NOT NULL  
, [SalesOrderNumber]     nvarchar(20) NOT NULL  
, [SalesOrderLineNumber] int      NULL  
, [SalesQuantity]        int      NOT NULL  
, [SalesAmount]          money    NOT NULL  
)
```

```
WITH
```

```
( CLUSTERED COLUMNSTORE INDEX  
  DISTRIBUTION = ROUND_ROBIN  
)  
;
```



```
CREATE TABLE [build].[FactOnlineSales]
```

```
(  
    [OnlineSalesKey]      int      NOT NULL  
, [DateKey]              datetime NOT NULL  
, [StoreKey]             int      NOT NULL  
, [ProductKey]           int      NOT NULL  
, [PromotionKey]         int      NOT NULL  
, [CurrencyKey]          int      NOT NULL  
, [CustomerKey]          int      NOT NULL  
, [SalesOrderNumber]     nvarchar(20) NOT NULL  
, [SalesOrderLineNumber] int      NULL  
, [SalesQuantity]        int      NOT NULL  
, [SalesAmount]          money    NOT NULL  
)
```

```
WITH
```

```
( CLUSTERED COLUMNSTORE INDEX  
  DISTRIBUTION = HASH([ProductKey])  
)  
;
```




Hash Distributed

```
CREATE TABLE ProductSales
WITH (DISTRIBUTION=HASH(AccountID))
AS...
```

ProductSales – Raw Data

AccountID	SalesAmt	...
47	\$1,234.36	...
36	\$2,345.47	...
14	\$3,456.58	...
25	\$4,567.69	...
48	\$5,678.70	...
37	\$6,789.81	...
...

Hash(30)




Round Robin Distributed

```
CREATE TABLE ProductSales
WITH (DISTRIBUTION = ROUND_ROBIN)
AS ...
```

ProductSales – Raw Data

AccountID	SalesAmt	...
47	\$1,234.36	...
36	\$2,345.47	...
14	\$3,456.58	...
25	\$4,567.69	...
48	\$5,678.70	...
37	\$6,789.81	...
42	\$1,632.25	...
42	\$4,453.21	...
52	\$7,892.81	...
91	\$9,549.64	...
66	\$2,498.14	...
23	\$3,145.99	...




T-SQL: Create a Replicated table

```
CREATE TABLE dbo.DimCustomer
(
    CustomerKey          int          NOT NULL
,   GeographyKey        int          NULL
,   CustomerAlternateKey nvarchar(15) NOT NULL
,   Title               nvarchar(8)  NULL
,   FirstName           nvarchar(50)  NULL
,   LastName            nvarchar(50)  NULL
,   BirthDate           date          NULL
,   Gender              nvarchar(1)  NULL
,   EmailAddress        nvarchar(50)  NULL
,   YearlyIncome        money        NULL
,   DateFirstPurchase   date          NULL
)
WITH
(
    CLUSTERED COLUMNSTORE INDEX
,   DISTRIBUTION = REPLICATED
)
```



Indexing



Indexing Choices

- Row Store : Clustered Index
- Column Store : Clustered Columnstore Index
- Heap : No index
- Non-Clustered Index

Indexing tables

```
CREATE TABLE [dbo].[DimStore]
```

```
(  
    [StoreKey]                int                NOT NULL  
, [GeographyKey]             int                NOT NULL  
, [StoreName]                 nvarchar(100)      NOT NULL  
, [StoreType]                 nvarchar(15)       NULL  
, [StoreDescription]          nvarchar(300)     NOT NULL  
, [Status]                    nvarchar(20)      NOT NULL  
, [OpenDate]                  datetime      NOT NULL  
, [CloseDate]                 datetime      NULL  
, [ETLLoadID]                 int                NULL  
, [LoadDate]                  datetime      NULL  
, [UpdateDate]                datetime      NULL  
)
```

```
WITH
```

```
( CLUSTERED INDEX([StoreKey])  
  DISTRIBUTION = ROUND_ROBIN  
)  
;
```

Row

```
CREATE TABLE [dbo].[FactOnlineSales]
```

```
(  
    [OnlineSalesKey]          int                NOT NULL  
, [DateKey]                  datetime          NOT NULL  
, [StoreKey]                 int                NOT NULL  
, [ProductKey]               int                NOT NULL  
, [PromotionKey]             int                NOT NULL  
, [CurrencyKey]              int                NOT NULL  
, [CustomerKey]              int                NOT NULL  
, [SalesOrderNumber]         nvarchar(20)      NOT NULL  
, [SalesOrderLineNumber]     int                NULL  
, [SalesQuantity]            int                NOT NULL  
, [SalesAmount]              money            NOT NULL  
)
```

```
WITH
```

```
( CLUSTERED COLUMNSTORE INDEX  
  DISTRIBUTION = HASH([ProductKey])  
)  
;
```


Column

Distribution

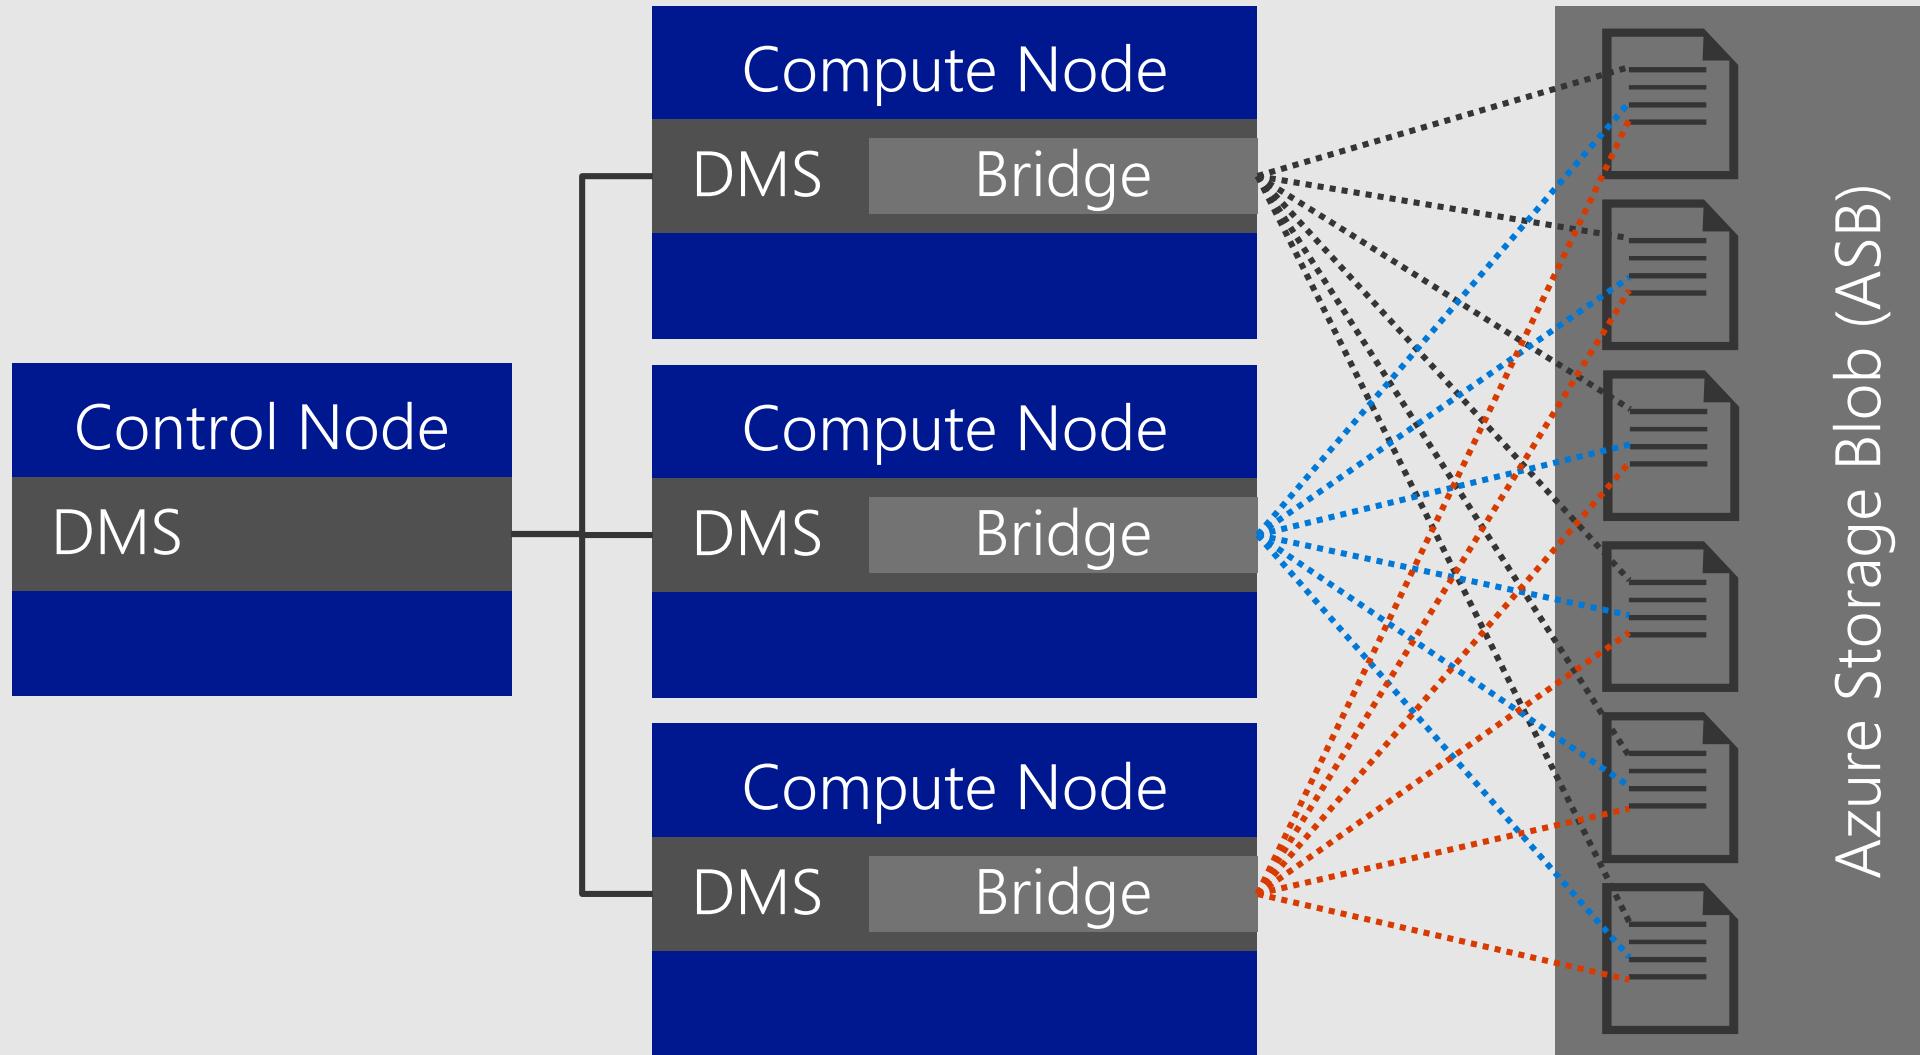
Clustered Columnstore & Partitioning

```
CREATE TABLE [dbo].[FactInternetSales]
(
    [ProductKey] int NOT NULL ,
    [OrderDateKey] int NOT NULL ,
    [CustomerKey] int NOT NULL ,
    [PromotionKey] int NOT NULL ,
    [SalesOrderNumber] nvarchar(20) NOT NULL ,
    [OrderQuantity] smallint NOT NULL ,
    [UnitPrice] money NOT NULL ,
    [SalesAmount] money NOT NULL
)
WITH ( CLUSTERED COLUMNSTORE INDEX ,
        DISTRIBUTION = HASH([ProductKey]) ,
        PARTITION ( [OrderDateKey] RANGE RIGHT FOR
VALUES (20000101,20010101,20020101,20030101,20040101,20050101)
        )
)
;
```

Polybase



Polybase parallel load from Azure Storage



Additional Resources

Azure site: <http://aka.ms/sqldw>

SQL DW suggestions: <http://aka.ms/sql-dw-feedback>

Stack Overflow – Tag: azure-sqldw

Twitter: @AzureSQLDW

Nominate: http://aka.ms/engage_dw_cse

Q&A



Modernizing **Your** Data Warehouse

