

Modernizing **Your** Data Warehouse






SQL DW Data Loading Best Practices

Casey Karst


cakarst@Microsoft.com



Agenda

- Table Structure and impacts to loading
- Loading patterns
- Loading tools
- Loading best practices

Physical Structures



Row store & Column store





Table and Index Terminology

Primary Data Stores

Heap = Base Row Store

Clustered Index (CI) = Base Row Store maintained as a B-Tree

Clustered Columnstore Index (CCI) = Base Column Store

Secondary Data Stores

Non Clustered Index (NCI) = Secondary B-Tree Index

- NCI can be on Heap

- NCI can be on Clustered Columnstore Index (NCI on CCI)

Non Clustered Columnstore Index (NCCI) = not supported in SQLDW

Column store taxonomy

Data Row Group Segments Column store




Row store or Column store?

- Small Data Set (< 60 million rows)
- Frequent updates
- Small Dimension tables




Row store



- Large Data Set (> 60 million rows)
- Mostly append only data
- Fact tables or large dimension tables



Column store



Why Partition?

Benefit to Loads

Data Lifecycle Management

- Drop partition avoids transaction logging

- Insert to empty table/partition avoids transaction logging

- ⇒ Partition Switching pattern

Targeted Index Builds

Benefit to Queries

Partition Elimination

Partitioning Guidance

Partition for data management

Lesser benefit had on partition elimination for faster performance

Don't over partition!


Partitioning granularity likely to differ to SQL Server

- Data is already spread across 60 distributions

Columnstore index row groups give ideal performance with 1 million rows each

Need at least 60 million per partition!

Loading patterns



Loading patterns

Pattern 1: Batch loads (PolyBase)

Load large volumes of data in parallel

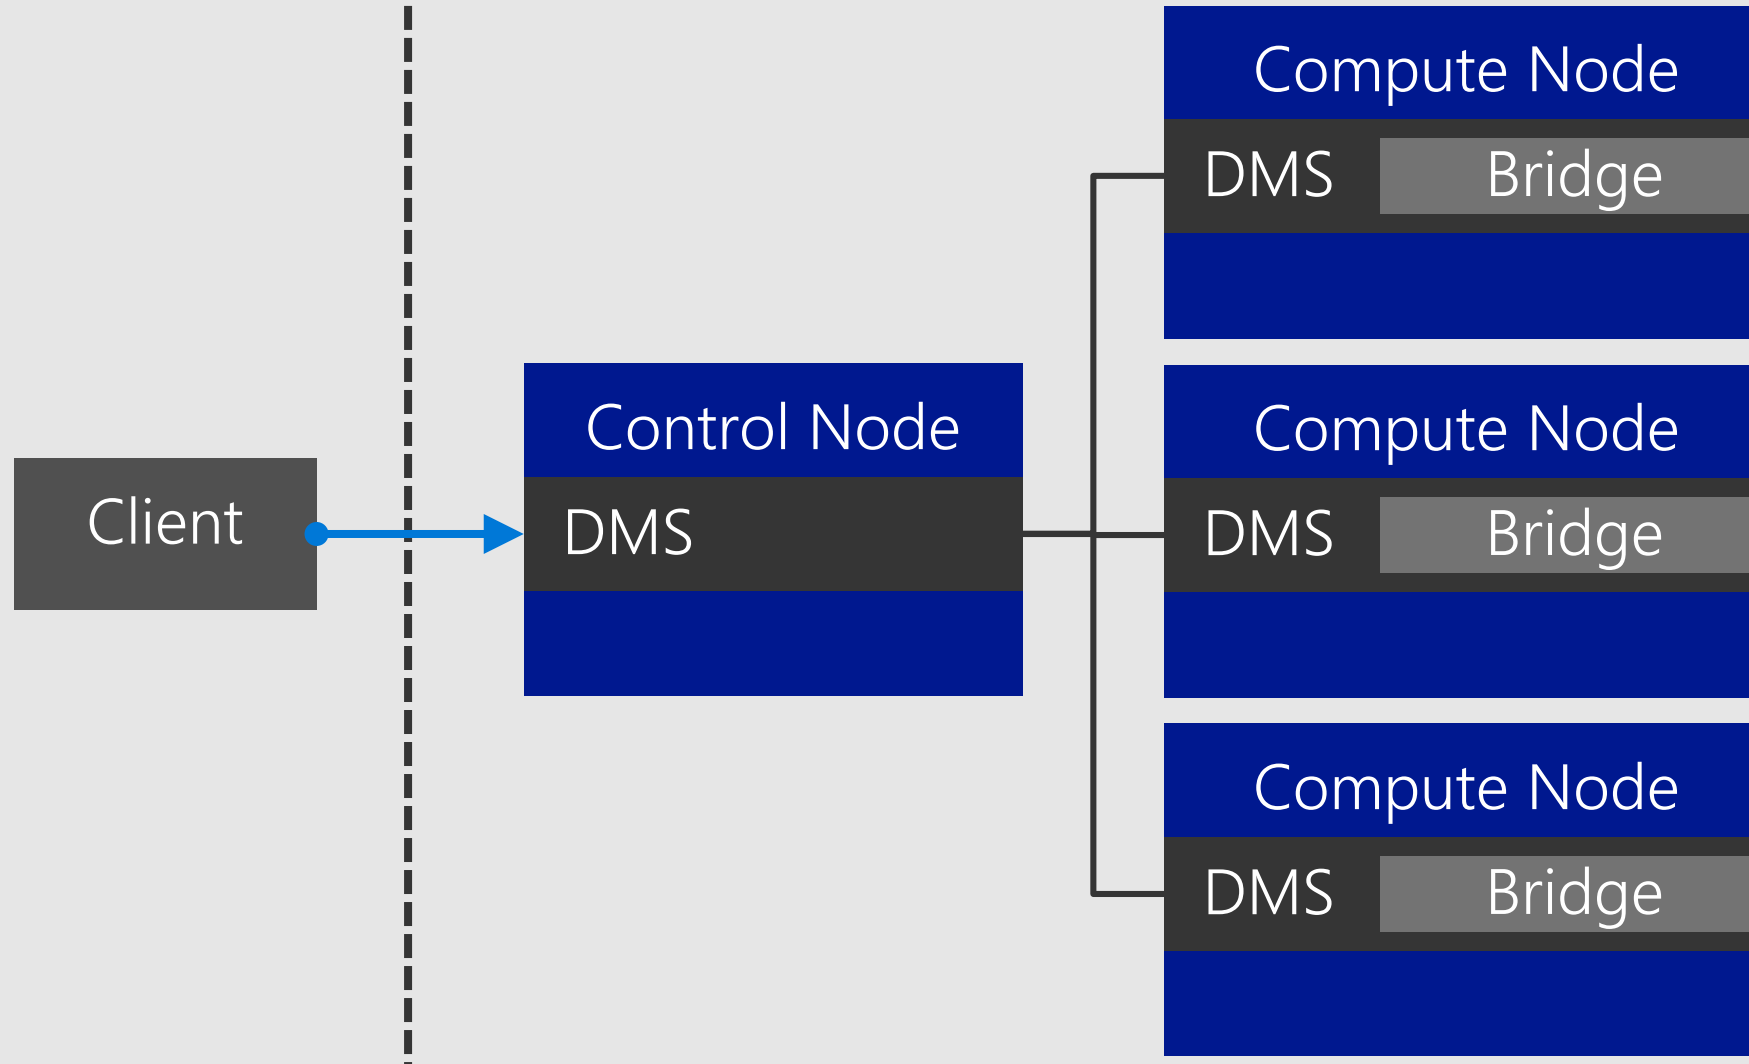
Move data from Azure storage to SQL DW in parallel to Compute nodes.

This includes Databricks integrations!

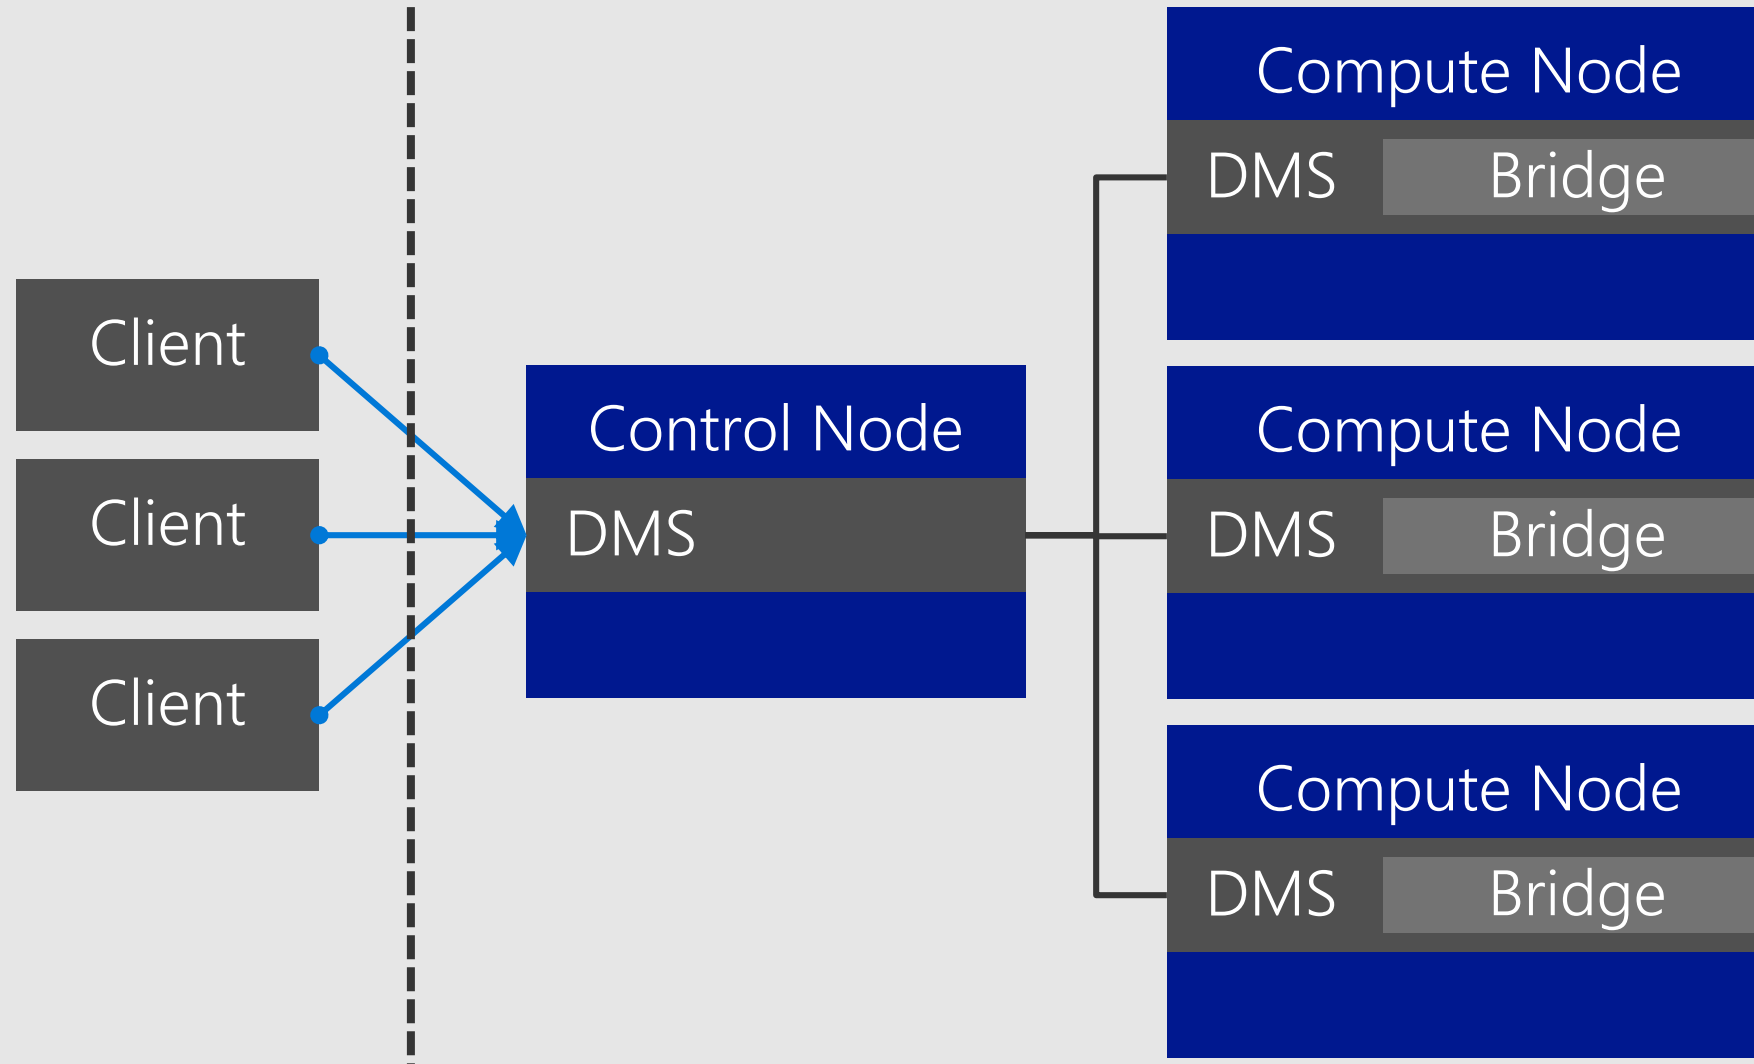
Pattern 2: Streaming loads (BCP)

- Single record or small batches in each load.
- Data moves from source to the Control node and then to the Compute nodes.

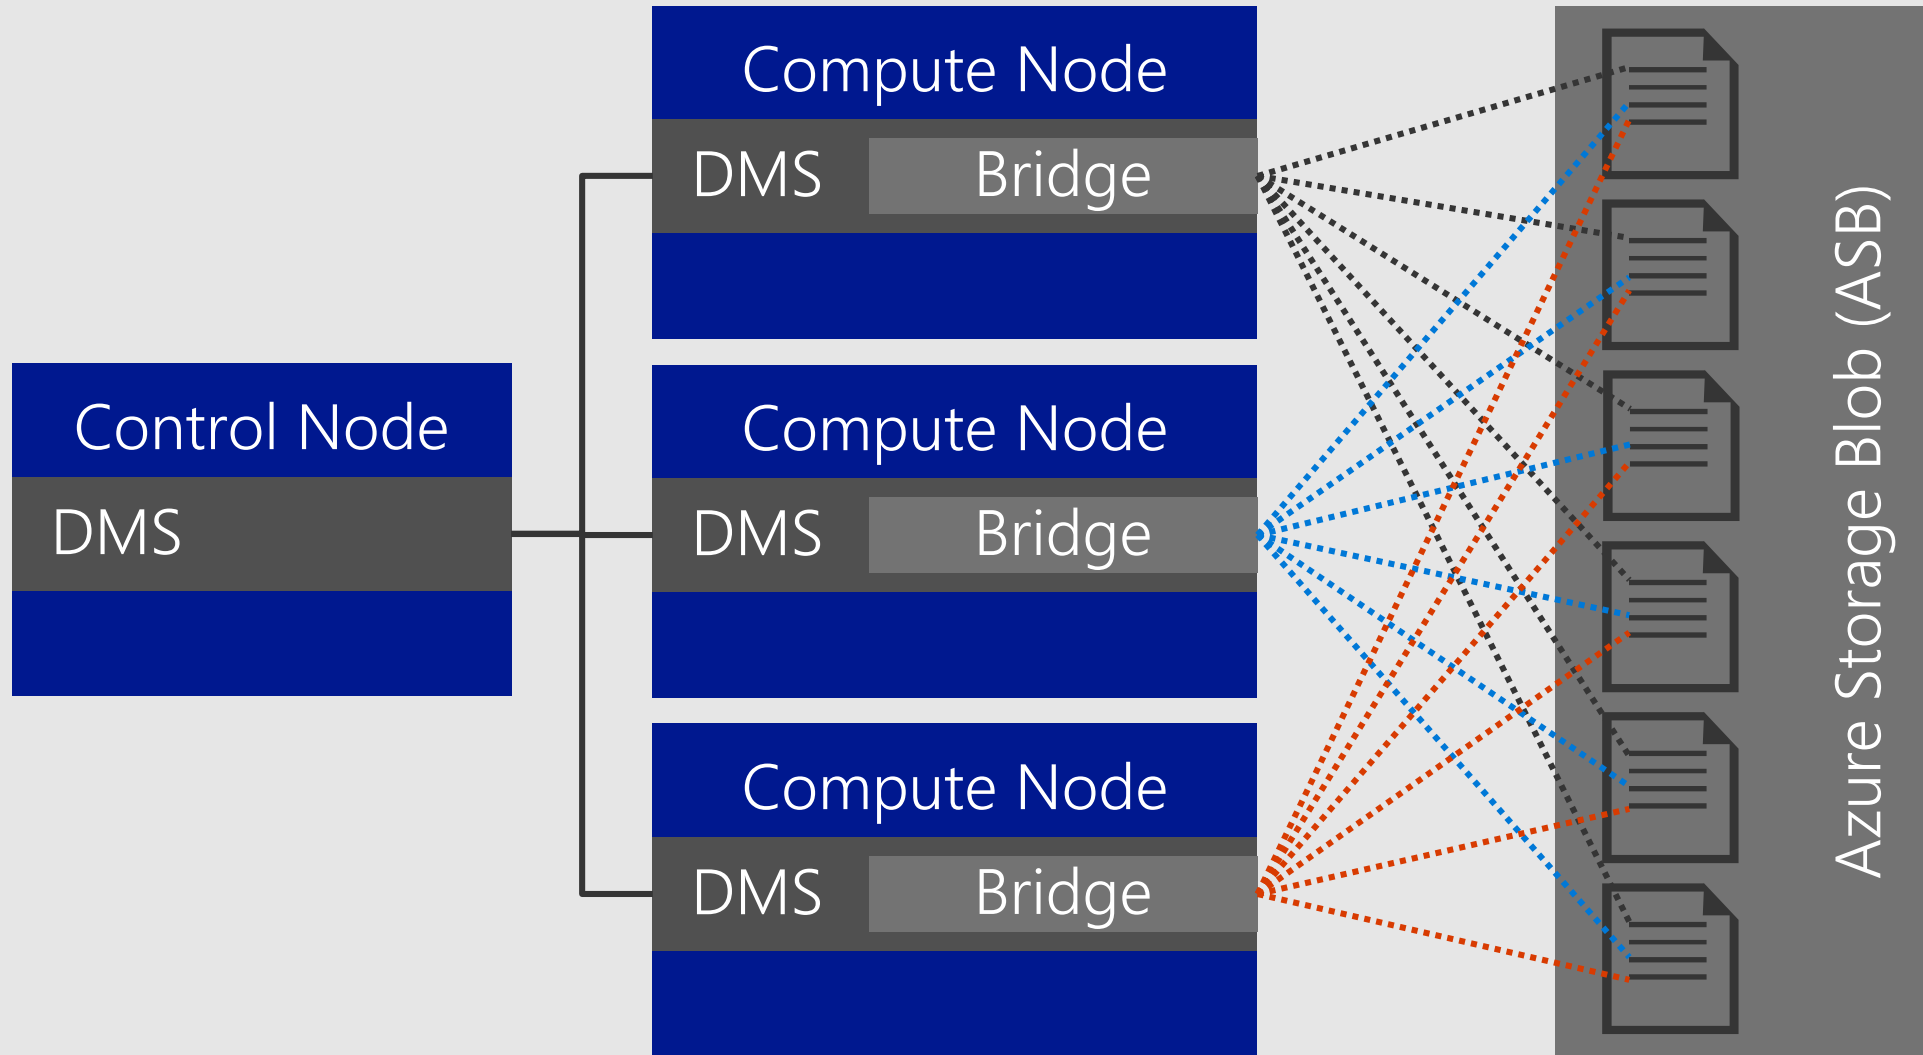
Single gated client




Single gated client parallelized



Polybase parallel load to Azure Storage Blob



Loading tools




Create external tables

```
CREATE EXTERNAL DATA SOURCE WASBStor
WITH (TYPE = Hadoop,
      LOCATION = 'wasbs://<container>@<account_name>.blob.core.windows.net',
      Credential = <Database scoped credential>);
```



Once per WASB container

```
CREATE EXTERNAL FILE FORMAT TextFile
WITH ( FORMAT_TYPE = DELIMITEDTEXT,
      DATA_COMPRESSION = 'org.apache.hadoop.io.compress.GzipCodec',
      FORMAT_OPTIONS (FIELD_TERMINATOR = '|', USE_TYPE_DEFAULT = TRUE));
```



Once per file format

```
CREATE EXTERNAL TABLE [dbo].[Customer_import] (
  [SensorKey] int NOT NULL,
  [CustomerKey] int NOT NULL,
  [Speed] float NOT NULL
)
WITH (LOCATION='<Dimensions/customer>',
      DATA_SOURCE = WASBStor,
      FILE_FORMAT = TextFile
)
```



File path

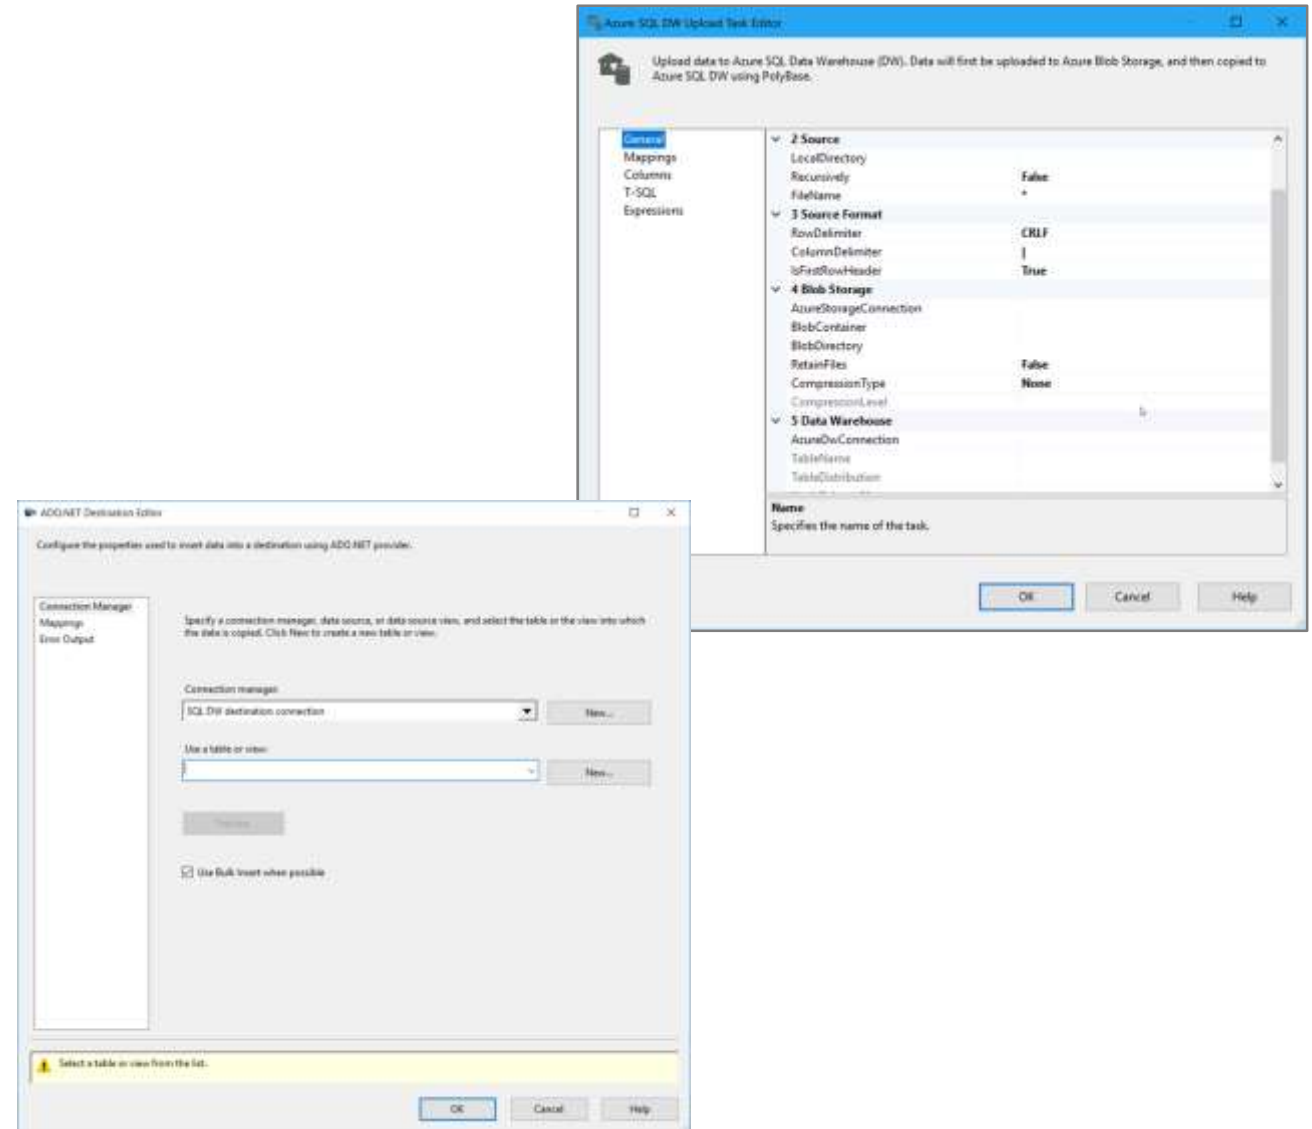
SQL Server Integration Services (SSIS)

Overview

SQL Server Integration Services is used to extract, transform data and load data from a variety of sources into Azure SQL Data Warehouse.

There are two options for loading data into SQL Data Warehouse with SSIS:

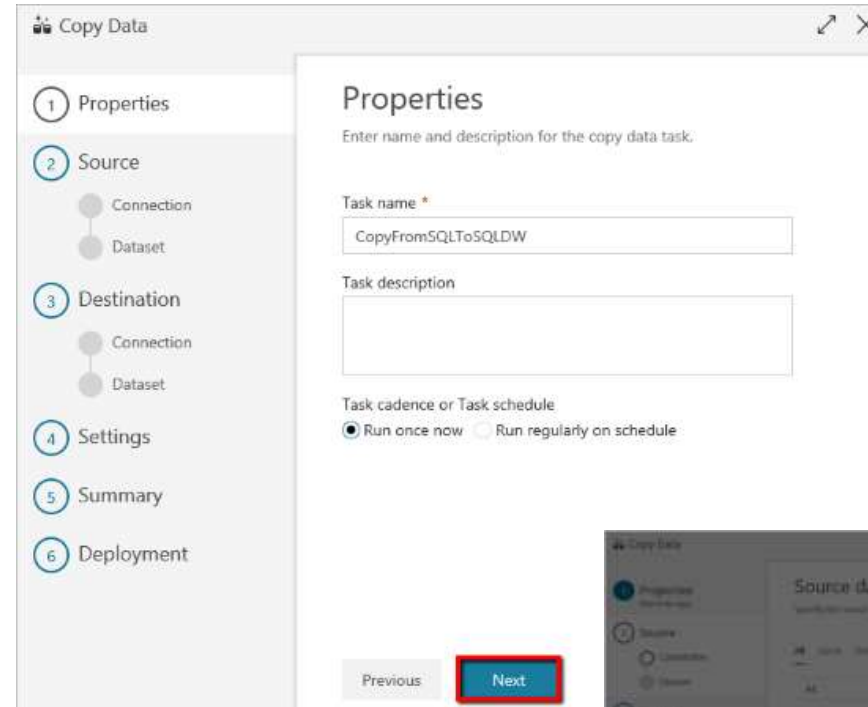
- **Azure SQL Data Warehouse Upload Task:** provides best performance but assumes source data is in delimited text file format.
- **Data Flow Task:** slower than SQL Data Warehouse Upload Task but supports a wider range of data sources.



Azure Data Factory Copy Data tool

Overview

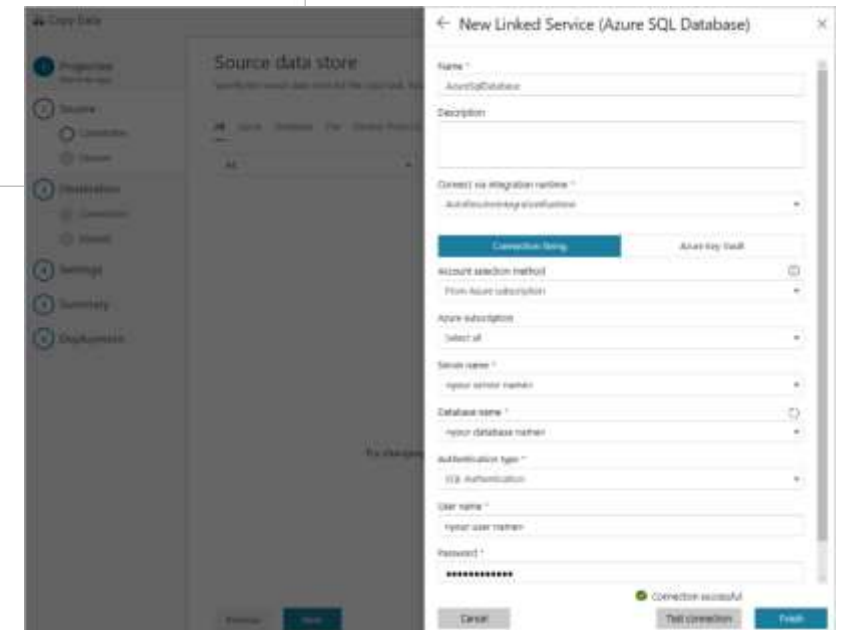
The Azure Data Factory Copy Data tool provides an intuitive wizard that allows you to copy data from a variety of data sources into Azure SQL Data Warehouse.



The screenshot shows the 'Copy Data' wizard in the 'Properties' step. On the left, a sidebar lists the steps: 1 Properties, 2 Source, 3 Destination, 4 Settings, 5 Summary, and 6 Deployment. The 'Properties' panel contains the following fields:

- Task name ***: A text box containing 'CopyFromSQLToSQLDW'.
- Task description**: An empty text box.
- Task cadence or Task schedule**: Two radio buttons, 'Run once now' (selected) and 'Run regularly on schedule'.

At the bottom of the panel are 'Previous' and 'Next' buttons. The 'Next' button is highlighted with a red rectangle.



The screenshot shows the 'New Linked Service (Azure SQL Database)' dialog. It contains the following fields and options:

- Name ***: A text box containing 'AzureSqlDatabase'.
- Description**: An empty text box.
- Connect via integration runtime ***: A dropdown menu with 'AzureSQLDWIntegrationRuntime' selected.
- Connection string** and **Azure key vault**: Two buttons, with 'Connection string' selected.
- Account selection method**: A dropdown menu with 'From Azure subscription' selected.
- Azure subscription**: A dropdown menu with 'Select all' selected.
- Server name ***: A text box containing 'myserver.mysql.database.azure.com'.
- Database name ***: A text box containing 'mydatabase'.
- Authentication type ***: A dropdown menu with 'SQL authentication' selected.
- User name ***: A text box containing 'myuser@myserver'.
- Password ***: A text box containing 'mypassword'.

At the bottom right, there is a green checkmark icon and the text 'Connection successful'. Below this are 'Cancel', 'Test connection', and 'Finish' buttons.

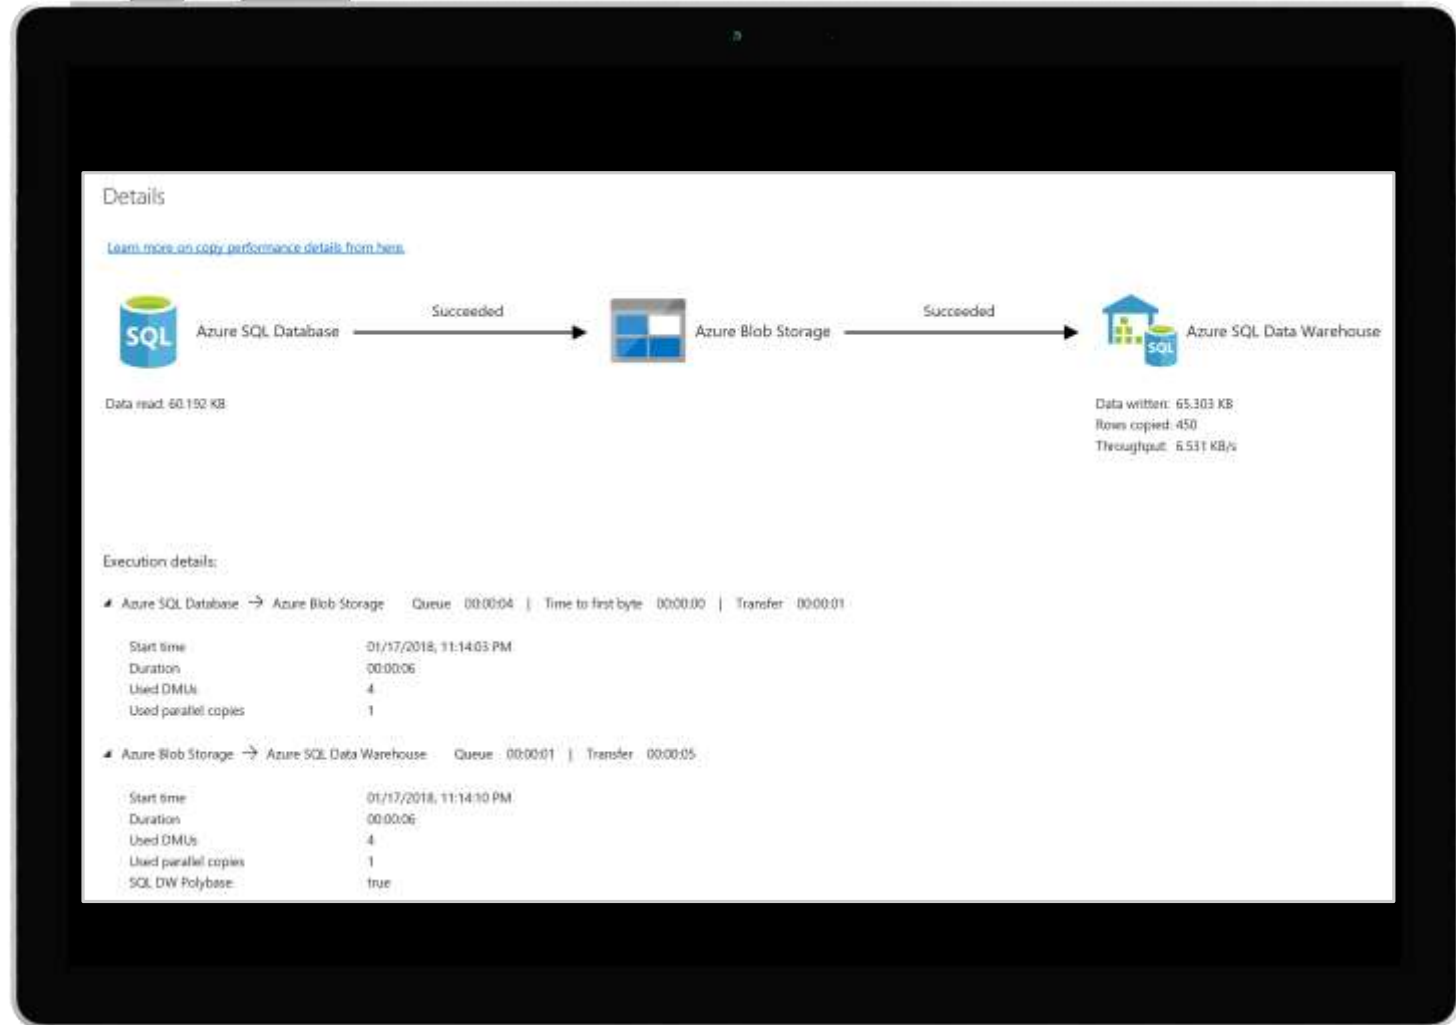
Azure Data Factory Copy activity

Overview

The Azure Data Factory Copy activity allows copying to and from Azure SQL Data Warehouse from any supported data store.

The Copy activity also supports retrieving data from a SQL source by using a SQL query or stored procedure. Authentication can be via:

- SQL Authentication
- Service principal token authentication
- Managed identity token authentication



Databricks – structured streaming

Overview

The Databricks SQL DW connector supports batch and structured streaming support for writing real-time data into Azure SQL Data Warehouse.

It uses Polybase and the Databricks structured streaming API to stream data from Kafka or Kinesis sources directly into SQL Data Warehouse at a user-configurable rate.

Source: <https://docs.azuredatabricks.net/spark/latest/data-sources/azure/sql-data-warehouse.html#streaming-support>

```
# Prepare streaming source; this could be Kafka,  
Kinesis, or a simple rate stream.
```

```
df = spark.readStream \  
    .format("rate") \  
    .option("rowsPerSecond", "100000") \  
    .option("numPartitions", "16") \  
    .load()
```

```
# Apply some transformations to the data then use  
# Structured Streaming API to continuously write the  
data to a table in SQL DW.
```

```
df.writeStream \  
    .format("com.databricks.spark.sqldw") \  
    .option("url", <azure-sqldw-jdbc-url>) \  
    .option("tempDir",  
"wasbs://<containername>@<storageaccount>.blob.core.  
windows.net/<directory>") \  
    .option("forwardSparkAzureStorageCredentials",  
"true") \  
    .option("dbTable", <table-name>) \  
    .option("checkpointLocation", "/tmp_location") \  
    .start()
```

Mechanism for loading

	PolyBase	SSIS	ADF	BCP	SqlBulkCopy
1. PolyBase 2. SSIS* 3. ADF 4. BCP 5. SQLBulkCopy API 6. Attunity Cloudbeam 7. ASA/Storm**	<div>Rate</div> <div>Fastest ← → Slowest</div>				
Rate increase as DWU increases	Yes	Yes	Yes	No	No
Rate increases as you add concurrent load	No	No	No	Yes	Yes


* With SSIS Azure Feature Pack June 2017 or newer

** Not a good idea

Loading method considerations

Loading Type	Source	Concerns	Advice
Batch loading	WASB/ADLS	Latency of data	Do it!
Micro Batch loading	WASB/ADLS	Potential impact to index health. Impact on machine resources	Do it with caution... Make sure that loads are big enough
Streaming Load	Azure Stream Analytics, BCP	Column store Index health. Load Performance	Do it with higher caution... Understand higher latency and impact on segment quality.

Best practices



Dimension tables - best practices

Use round robin or replicated for small tables

Use clustered index, not clustered columnstore index

Can load directly to production because of small size

Use metadata rename to reload data

Fact table best practices

Use partitions to reduce the loading impact on the production table

Consider landing data from ADL in a staging table

Take advantage of directory structure to limit loading scope

More reading...

- Guidance for designing distributed tables
- <https://docs.microsoft.com/en-us/azure/sql-data-warehouse/sql-data-warehouse-tables-distribute>
- Columnstore indexes
- <https://docs.microsoft.com/en-us/sql/relational-databases/indexes/columnstore-indexes-overview?view=sql-server-2017>
- Analyze your workload in Azure SQL Data Warehouse
- <https://docs.microsoft.com/en-us/azure/sql-data-warehouse/analyze-your-workload>
- Adaptive caching powers Azure SQL Data Warehouse performance gains
- <https://azure.microsoft.com/en-us/blog/adaptive-caching-powers-azure-sql-data-warehouse-performance-gains/>
- Cheat sheet
- <https://docs.microsoft.com/en-us/azure/sql-data-warehouse/cheat-sheet>

Q&A

